# Tactical Adaptation or Good Fortune: A Spatial Analysis of Bayer Leverkusen's Late-Game Pressing

Barcelona School of Economics | Geospatial Data Science | Term Paper

Daniel Campos, Nikoloz Darsalia, Tina Sikharulidze, and Eric Gutiérrez

2026-03-27

## Contents

**Introduction** **1**

**Setup and Libraries** **3**

**Data Collection** **3**
StatsBomb: Match Events . . . 3
ClubElo: Opponent Quality Ratings . . . 4

**Descriptive Analysis of Bayer Leverkusen's Spatial Patterns** **7**
Building the Zonal Counterpress Rate Dataset . . . 7
Building the Zonal Counterpress Raster . . . 8
KDE Heatmaps: Press Density vs Counterpress Density . . . 10
Raster Analysis: Trap Rate Surface and Zonal Statistics . . . 13

**The Comeback Shift: A DiD Approach** **17**
Score State Reconstruction . . . 17
Building the DiD Dataset . . . 18
DiD: Regression Models . . . 20
DiD: Visualizations . . . 22

**Extension: Benchmarking Against 2015/16 Bundesliga Top Teams** **41**

**Conclusion** **53**

**References** **54**

---

## Introduction

**Motivation:**

On May 18, 2024, Bayer Leverkusen finished the most remarkable season in Bundesliga history: 34 matches unbeaten, including ten comebacks from losing positions. A team that was trailing in the final minutes should have dropped points many times, yet they never did. Most analysts and journalists put this down to resilience, mentality, or luck, but Xabi Alonso was clear that it came from their tactical system. That difference is important: if a team's ability to come back when losing is something structural, built into how they press and move, rather than something that just happens, it changes how we think about what makes

a winning team. This study aims to test that idea, shifting the focus of this season away from subjective explanations and toward something measurable. If their unbeaten identity was truly a system, it should leave a clear trace on the pitch.

Whether late-game performance reflects deliberate tactics or good fortune is something football science has not yet answered in a rigorous way. The field has advanced considerably in recent years, using spatio-temporal tracking data to study pitch control, passing networks, and how game state shapes team behavior, showing for example how teams that go ahead tend to drop their defensive line and give up possession. Pressing behavior has been studied spatially (Andrienko et al., 2017), automated tools have been developed to detect and measure counterpressing across full seasons (Bauer & Anzer, 2021; Peters et al., 2025), the effect of game state on ball recovery timing has been quantified (Vogelbein et al., 2014), and causal inference has been brought into football research through set-piece analysis like throw-ins (Epasinghege Dona & Swartz, 2023). To the best of our knowledge, what has not been done is combining all of this to ask whether a specific team pushes its pressing line higher in a systematic way when losing late in a match, tested with formal causal inference on spatial data. Where a team presses under pressure reflects what has been trained on the training ground, and Leverkusen's unusual comeback record makes it the best possible case to study: if a spatial signal exists anywhere, it should appear here.

The analysis comes with real challenges. The main problem is partly caused by the team itself: because Leverkusen were so rarely losing late in matches, the number of situations available to study is small, which limits what traditional statistical methods can confidently pick up. Spatial event data also adds difficulty, since coordinated pressing and individual desperate running can look similar in the numbers. These are not easy problems, but they are not impossible ones either. A pressing system that was genuinely designed to handle any deficit should produce a spatial pattern clear enough to detect even in noisy data. This brings us to our main research question: does Leverkusen's 2023/24 season reflect a distinctive tactical system, specifically an unusual spatial pressing response when losing, or is aggressive late-game pressing simply something all top Bundesliga teams do?

**Preview of the Results:**

We find that it does. Leverkusen presses 6.2 meters higher when losing after the 70th minute compared to the counterfactual trend, a gap roughly equal to the width of the penalty area. However, the results obtained are not statistically significant.

Against their own 2015/16 incarnation, the contrast is even more defined: the same club, nine years earlier, dropped their press when losing late by 3.7 meters. The comeback identity, in other words, is not the club; it is Alonso's system.

To arrive at these findings we combine three analytical layers. First, we map where Leverkusen counter-presses using kernel density estimation and raster zonal statistics, finding that wide attacking channels are systematically exploited as pressing traps. Second, we use a Difference-in-Differences design to isolate the causal effect of a losing score state on press height in the final 20 minutes of matches. Third, we benchmark Leverkusen 2023/24 against the top seven teams from the 2015/16 Bundesliga season, showing that their pressing profile in losing-late conditions sits in a quadrant no comparable team occupies: higher absolute press height and stronger maintenance of that height under pressure simultaneously.

The analysis is made possible by StatsBomb's open event data, which records the precise spatial coordinates of every pressing action across the full season.

**Paper Structure:**

The structure of the document is as follows. First, in section 2, we list the required packages for the project. In section 3, the data collection process is detailed, which consists on retrieving StatsBomb event data and ClubElo opponent ratings. Then, in section 4, we perform some descriptive analysis of the data, finding where does Bayer Leverkusen press and counterpress. Finally, in section 5, we perform a Difference-in-Differences to verify whether Bayer Leverkusen systematically shifts their press forward when losing late. Additionally, we also include an extension that benchmarks the team's performance against 2015/16 Bundesliga top teams. Section 6 concludes.

---

## Setup and Libraries

First, we load the necessary R packages to handle the complete pipeline of our analysis. We rely on libraries like `dplyr` and `tidyr` for robust data manipulation and cleaning. To access our raw event data, we use `StatsBombR`, which connects directly to StatsBomb's free API. Because our research question centers heavily on tactical geometry, we incorporate specialized spatial libraries: `sf` for managing vector data (such as individual pressing events) and `terra` for raster operations (like continuous density heatmaps). Finally, we load our visualization toolkit, featuring `ggplot2` and `ggsoccer`, to accurately map our spatial findings onto a standardized football pitch.

```r
# Core data manipulation
library(dplyr)
library(tidyr)
library(lubridate)
library(readr)
library(broom)

# StatsBomb free data
library(StatsBombR)

# Vector spatial data
library(sf)

# Raster spatial data
library(terra)

# Visualization
library(ggplot2)
library(ggsoccer)   # football pitch annotation layer
library(patchwork)  # combine ggplot panels
library(scales)
```

---

## Data Collection

### StatsBomb: Match Events

We use the StatsBomb open data API via the `StatsBombR` package. The data covers all Bayer Leverkusen matches in the 2023/24 Bundesliga season.

Notice that StatsBomb uses a flat 2D pitch with x  [0, 120] and y  [0, 80] as coordinate system. Leverkusen always attacks towards x = 120. All spatial objects in this file use this coordinate system with no geographic CRS (`crs = NA_crs_`).

```r
# Download all matches for the 2023/24 Bundesliga season (competition_id = 9)
invisible(capture.output({
  comps <- FreeCompetitions() %>%
    filter(competition_id == 9, season_name == "2023/2024")

  matches <- FreeMatches(comps)
}))
```

```r
# Keep only Leverkusen matches
leverkusen_matches <- matches %>%
  filter(
    home_team.home_team_name == "Bayer Leverkusen" |
    away_team.away_team_name == "Bayer Leverkusen"
  )

cat("Leverkusen matches found:", nrow(leverkusen_matches), "\n")
```

```
## Leverkusen matches found: 34
```

```r
# Download all events for those matches.
# allclean() parses nested JSON columns (location, shot, etc.) into flat columns.
cat("Downloading full-season event data...\n")
```

```
## Downloading full-season event data...
```

```r
invisible(capture.output({
events_raw <- free_allevents(MatchesDF = leverkusen_matches)
events <- allclean(events_raw)
}))

# Filter to Leverkusen pressure events
pressure_events <- events %>%
  filter(
    team.name == "Bayer Leverkusen",
    type.name == "Pressure",
    !is.na(location.x),
    !is.na(location.y)
  )

cat("Total Leverkusen pressure events:", nrow(pressure_events), "\n")
```

```
## Total Leverkusen pressure events: 4786
```

## ClubElo: Opponent Quality Ratings

We fetch ClubElo ratings, an estimation of its strength based on past results, for each opponent on the day of the match. This allows us to stratify the analysis by opponent quality tier.

```r
# Helper function: returns the ClubElo rating for a team on a given date.
# The API returns a historical table; we find the row whose date range contains
# the match date, with a fallback to the most recent rating before the date.
get_elo_on_date <- function(team, date) {
  date <- as.Date(date)
  url <- paste0("http://api.clubelo.com/", team)
  hist <- suppressWarnings(read_csv(url, show_col_types = FALSE)) %>%
    mutate(From = as.Date(From), To = as.Date(To)) %>%
    arrange(From)

  # Find a row where date falls within [From, To]
  in_range <- hist %>%
    filter(!is.na(From) & date >= From & (is.na(To) | date <= To))
  if (nrow(in_range) >= 1) return(as.numeric(in_range$Elo[1]))

  # Fallback: use the most recent rating before the match date
```

```r
  before <- hist %>% filter(!is.na(From) & From <= date)
  if (nrow(before) >= 1) return(as.numeric(before$Elo[nrow(before)]))

  NA_real_
}
```

```r
# For simplicity, we compute fixed ELO tier thresholds from a pre-season anchor
# date. This ensures thresholds do not shift mid-season as ratings update.
anchor_date <- as.Date("2023-08-18")

bundesliga_anchor <- read_csv(
  paste0("http://api.clubelo.com/", format(anchor_date, "%Y-%m-%d")),
  show_col_types = FALSE
) %>%
  filter(
    Country == "GER",
    Level == 1,
    as.Date(From) <= anchor_date,
    as.Date(To) >= anchor_date
  )

# Tiers: bottom third = Weak, middle = Mid, top third = Top
elo_cuts <- quantile(bundesliga_anchor$Elo, probs = c(1/3, 2/3), type = 7,
                     na.rm = TRUE)
p33 <- unname(elo_cuts[1])
p67 <- unname(elo_cuts[2])

cat(sprintf("ELO thresholds (anchor: %s)\n  P33 = %.0f  |  P67 = %.0f\n",
            anchor_date, p33, p67))
```

```
## ELO thresholds (anchor: 2023-08-18)
##   P33 = 1636  |  P67 = 1722
```

```r
# Build a match-level table with opponent name, match date, and ELO rating.
# StatsBomb team names differ from ClubElo records, we standardize them.
match_info <- leverkusen_matches %>%
  mutate(
    opponent = ifelse(
      home_team.home_team_name == "Bayer Leverkusen",
      away_team.away_team_name,
      home_team.home_team_name
    ),
    match_date = as.Date(match_date),
    opponent_clubelo = case_when(
      opponent == "Bayern Munich"             ~ "Bayern",
      opponent == "Borussia Dortmund"         ~ "Dortmund",
      opponent == "RB Leipzig"                ~ "RBLeipzig",
      opponent == "VfB Stuttgart"             ~ "Stuttgart",
      opponent == "Eintracht Frankfurt"       ~ "Frankfurt",
      opponent == "Borussia Mönchengladbach"  ~ "Gladbach",
      opponent == "VfL Wolfsburg"             ~ "Wolfsburg",
      opponent == "SC Freiburg"               ~ "Freiburg",
      opponent == "FC Augsburg"               ~ "Augsburg",
      opponent == "TSG Hoffenheim"            ~ "Hoffenheim",
      opponent == "Werder Bremen"             ~ "Werder",
```

```r
      opponent == "FC Heidenheim"             ~ "Heidenheim",
      opponent == "Darmstadt 98"              ~ "Darmstadt",
      opponent == "FC Köln"                   ~ "Koeln",
      opponent == "FSV Mainz 05"              ~ "Mainz",
      opponent == "VfL Bochum 1848"           ~ "Bochum",
      TRUE ~ gsub(" ", "", opponent)
    )
  ) %>%
  select(match_id, match_date, opponent, opponent_clubelo)

# Fetch opponent ELO on match date (one API call per match)
cat("Fetching opponent ELO ratings...\n")
```

```
## Fetching opponent ELO ratings...
```

```r
match_info <- match_info %>%
  rowwise() %>%
  mutate(opponent_elo = get_elo_on_date(opponent_clubelo, match_date)) %>%
  ungroup() %>%
  mutate(
    elo_tier = case_when(
      is.na(opponent_elo)    ~ "Unknown",
      opponent_elo < p33     ~ "Weak opponent",
      opponent_elo < p67     ~ "Mid opponent",
      TRUE                   ~ "Top opponent"
    ),
    elo_tier = factor(elo_tier,
                      levels = c("Top opponent", "Mid opponent", "Weak opponent"))
  )

# Join ELO information onto every pressure event
pressure_events_final <- pressure_events %>%
  left_join(match_info, by = "match_id") %>%
  mutate(
    elo_tier = case_when(
      is.na(opponent_elo)    ~ "Unknown",
      opponent_elo < p33     ~ "Weak opponent",
      opponent_elo < p67     ~ "Mid opponent",
      TRUE                   ~ "Top opponent"
    ),
    elo_tier = factor(elo_tier,
                      levels = c("Top opponent", "Mid opponent", "Weak opponent"))
  )

cat("Pressure events with ELO assigned:", nrow(pressure_events_final), "\n")
```

```
## Pressure events with ELO assigned: 4786
```

```r
cat("ELO tier distribution:\n")
```

```
## ELO tier distribution:
```

```r
print(table(pressure_events_final$elo_tier))
```

```
##
##  Top opponent  Mid opponent Weak opponent
```

```
##          1203          1924          1659
```

---

## Descriptive Analysis of Bayer Leverkusen's Spatial Patterns

Before asking when the press shifts, we establish where Bayern Leverkusen systematically try to force turnovers. This section combines KDE heatmaps, vector zone polygons, and raster zonal statistics to characterize the spatial structure of pressing traps.

### Building the Zonal Counterpress Rate Dataset

```r
# In StatsBomb, a counterpress event (pressing immediately after losing the
# ball) has counterpress = TRUE. NA means it was a regular press, so we recode
# NA as FALSE.
press_clean <- pressure_events_final %>%
  filter(!is.na(location.x), !is.na(location.y)) %>%
  mutate(
    counterpress_flag = ifelse(is.na(counterpress), FALSE, TRUE),

    # Assign each press to one of 7 spatial zones:
    # Three pitch thirds (Defensive / Middle / Attacking) ×
    # Three lateral channels (Left / Central / Right), collapsed at D3
    zone6 = case_when(
      location.x <  40                              ~ "D3",
      location.x <  80 & location.y <  26.7         ~ "M3_L",
      location.x <  80 & location.y <  53.3         ~ "M3_C",
      location.x <  80                              ~ "M3_R",
      location.x >= 80 & location.y <  26.7         ~ "A3_L",
      location.x >= 80 & location.y <  53.3         ~ "A3_C",
      TRUE                                          ~ "A3_R"
    ),
    zone6 = factor(zone6,
                   levels = c("D3", "M3_L", "M3_C", "M3_R",
                              "A3_L", "A3_C", "A3_R"))
  )

cat("--- Pressing Trap Summary ---\n")
```

```
## --- Pressing Trap Summary ---
```

```r
cat("Total press events:", nrow(press_clean), "\n")
```

```
## Total press events: 4786
```

```r
cat("Counterpress events:", sum(press_clean$counterpress_flag), "\n")
```

```
## Counterpress events: 1098
```

```r
cat("Baseline CP rate:",
    round(mean(press_clean$counterpress_flag) * 100, 1), "%\n\n")
```

```
## Baseline CP rate: 22.9 %
```

```r
# Counterpress rate by spatial zone
press_clean %>%
```

```r
  group_by(zone6) %>%
  summarise(
    n_presses = n(),
    n_counterpress = sum(counterpress_flag),
    cp_rate = round(mean(counterpress_flag) * 100, 1),
    .groups = "drop"
  ) %>%
  arrange(desc(cp_rate))
```

```
## # A tibble: 7 x 4
##   zone6 n_presses n_counterpress cp_rate
##   <fct>     <int>          <int>   <dbl>
## 1 A3_R        537            193    35.9
## 2 A3_L        449            159    35.4
## 3 A3_C        641            182    28.4
## 4 M3_C        543            132    24.3
## 5 M3_L        683            158    23.1
## 6 M3_R        732            164    22.4
## 7 D3         1201            110     9.2
```

### Building the Zonal Counterpress Raster

We formally build the 7 pitch zones as `sf` polygon objects. This allows us to use the `sf` spatial join toolkit (`st_join`, `st_intersects`) to assign zone membership to each press event.

```r
# Helper function: creates a rectangular sf POLYGON from coordinate bounds.
# Returns an sf data frame with a zone6 label column and flat CRS.
make_zone_polygon <- function(xmin, xmax, ymin, ymax, name) {
  # Build the ring of 5 coordinates (close the polygon by repeating first point)
  coords <- matrix(
    c(xmin, ymin,
      xmax, ymin,
      xmax, ymax,
      xmin, ymax,
      xmin, ymin),
    ncol = 2, byrow = TRUE
  )
  st_sf(
    zone6 = name,
    geometry = st_sfc(st_polygon(list(coords)), crs = NA_crs_)
  )
}

# Create all 7 zone polygons and bind into a single sf object
zone_polygons <- bind_rows(
  make_zone_polygon(0, 40, 0, 80, "D3"),
  make_zone_polygon(40, 80, 0, 26.7, "M3_L"),
  make_zone_polygon(40, 80, 26.7, 53.3, "M3_C"),
  make_zone_polygon(40, 80, 53.3, 80, "M3_R"),
  make_zone_polygon(80, 120, 0, 26.7, "A3_L"),
  make_zone_polygon(80, 120, 26.7, 53.3, "A3_C"),
  make_zone_polygon(80, 120, 53.3, 80, "A3_R")
)
```

```r
cat("Zone polygons created:\n")
```

```
## Zone polygons created:
```

```r
print(zone_polygons %>% st_drop_geometry())
```

```
##   zone6
## 1    D3
## 2  M3_L
## 3  M3_C
## 4  M3_R
## 5  A3_L
## 6  A3_C
## 7  A3_R
```

```r
# Convert press events to sf POINT objects
press_sf <- press_clean %>%
  st_as_sf(
    coords = c("location.x", "location.y"),
    crs = NA_crs_,
    remove = FALSE
  )

# Spatial join: assign zone to each press point
# st_intersects checks which zone polygon contains each press point
press_sf_joined <- press_sf %>%
  st_join(zone_polygons, join = st_intersects, suffix = c("_orig", "_spatial"))

cat("Events successfully assigned via spatial join:",
    sum(!is.na(press_sf_joined$zone6_spatial)), "\n")
```

```
## Events successfully assigned via spatial join: 4805
```

```r
cat("Events on zone boundary (NA):",
    sum( is.na(press_sf_joined$zone6_spatial)), "\n")
```

```
## Events on zone boundary (NA): 0
```

```r
# Compute zone polygon centroids for map label placement
# st_centroid returns the geometric center of mass of each polygon
zone_centroids <- zone_polygons %>%
  st_centroid() %>%
  mutate(
    cx = st_coordinates(.)[, 1],   # extract x from sfc geometry
    cy = st_coordinates(.)[, 2]    # extract y from sfc geometry
  ) %>%
  st_drop_geometry()  # drop geometry column

print(zone_centroids)
```

```
##   zone6  cx    cy
## 1    D3  20 40.00
## 2  M3_L  60 13.35
## 3  M3_C  60 40.00
## 4  M3_R  60 66.65
## 5  A3_L 100 13.35
## 6  A3_C 100 40.00
```

```
## 7  A3_R 100 66.65
```

### KDE Heatmaps: Press Density vs Counterpress Density

The KDE maps show the spatial density of (A) all pressing events and (B) only counterpressing events. A counterpressing event is usually defined as an attempt to win back possession upon losing it, and includes the following: pressure, dribbled past, 50-50, duel, block, interception, and foul committed (not offensive). In our case, the time threshold to classify any of these actions as counterpressing events is set at 5 seconds.

```r
# Helper function: build a KDE heatmap on a StatsBomb pitch
make_kde_map <- function(data, title, subtitle,
                         fill_palette = "magma", legend_label = "Density") {
  ggplot(data) +
    stat_density_2d(
      aes(x = location.x, y = location.y, fill = after_stat(ndensity)),
      geom = "raster",
      contour = FALSE,
      n = 200,
      h = c(12, 12)   # bandwidth in pitch meters
    ) +
    scale_fill_viridis_c(
      option = fill_palette,
      direction = -1,
      begin = 0.05,
      name = legend_label,
      limits = c(0, 1)
    ) +
    # Mask rendering outside pitch boundaries with white rectangles
    annotate("rect", xmin=-Inf, xmax=Inf, ymin=-Inf, ymax=0, fill="white",
             color=NA) +
    annotate("rect", xmin=-Inf, xmax=Inf, ymin=80, ymax=Inf, fill="white",
             color=NA) +
    annotate("rect", xmin=-Inf, xmax=0, ymin=-Inf, ymax=Inf, fill="white",
             color=NA) +
    annotate("rect", xmin=120, xmax=Inf, ymin=-Inf, ymax=Inf, fill="white",
             color=NA) +
    annotate_pitch(dimensions = pitch_statsbomb, colour = "grey30", fill = NA) +
    # Zone grid reference lines
    geom_vline(xintercept = c(40, 80), colour = "white",
               linewidth = 0.5, linetype = "dashed", alpha = 0.6) +
    geom_hline(yintercept = c(26.7, 53.3), colour = "white",
               linewidth = 0.5, linetype = "dashed", alpha = 0.6) +
    coord_fixed(xlim = c(-2, 122), ylim = c(-5, 83), expand = FALSE) +
    labs(title = title, subtitle = subtitle) +
    theme_pitch() +
    theme(
      plot.title = element_text(face = "bold", size = 13, hjust = 0.5),
      plot.subtitle = element_text(size = 10, hjust = 0.5, colour = "grey40"),
      legend.position = "right",
      plot.background = element_rect(fill = "white", color = NA),
      panel.background = element_rect(fill = "white", color = NA)
    )
}

# Map A: all pressing events
```

```r
map_all <- make_kde_map(
  press_clean,
  title = "A) All Pressing Events",
  subtitle = "Spatial density: where Leverkusen presses",
  fill_palette = "Blues",
  legend_label = "Pressing\nDensity"
)

# Map B: counterpressing events only
map_cp <- make_kde_map(
  press_clean %>% filter(counterpress_flag == TRUE),
  title = "B) Counterpressing Events",
  subtitle = "Spatial density: where Leverkusen counterpresses",
  fill_palette = "magma",
  legend_label = "Trap\nDensity"
)

fig1 <- map_all + map_cp +
  plot_annotation(
    title = "Bayer Leverkusen 2023/24: Pressing Density vs Trap Density",
    theme = theme(
      plot.title = element_text(face = "bold", size = 15, hjust = 0.5),
      plot.subtitle = element_text(size = 11, hjust = 0.5, colour = "grey40")
    )
  )

print(fig1)
```

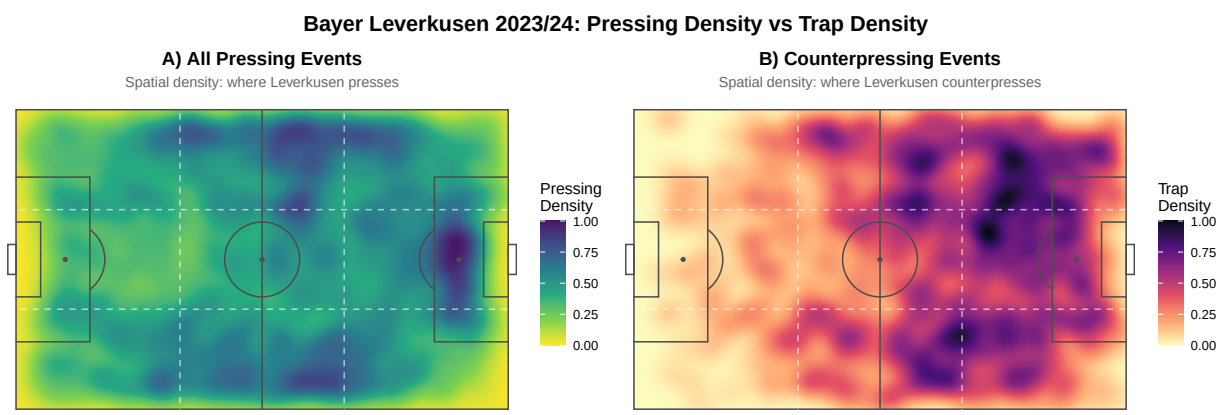


```r
ggsave("fig1_press_vs_trap_density.png", plot = fig1,
       width = 14, height = 6, dpi = 300)
```

The left panel (A) visualizes the spatial density of Bayer Leverkusen's total pressing events throughout the 2023/24 season. The heatmap reveals a striking concentration of pressing activity deep in the opponent's penalty area, represented by the distinct dark blue tone on the right side of the pitch. Beyond this intense focal point, the team's general pressing footprint is moderately distributed across the middle third and flanks, indicating a structured approach to applying pressure outside of the final third.

The right panel (B) shifts the focus to counterpressing events, illustrating exactly where Leverkusen immediately presses right after losing the ball in an attempt to regain possesion. Unlike the highly localized spike seen in their overall pressing, the density of these counterpressing events is broadly distributed across the entire attacking half. The heavy clustering of high-density dark purple zones, which extends from the center circle through to the opponent's final third, demonstrates that their tactical system prioritizes not allowing the opponent to build a potentially effective attacking play by pressing up high, rather than just recovering a balanced structure to face the defensive phase.

```r
# Repeat for each ELO tier
tier_maps <- lapply(
  c("Top opponent", "Mid opponent", "Weak opponent"),
  function(tier) {
    d <- press_clean %>% filter(counterpress_flag == TRUE, elo_tier == tier)
    make_kde_map(
      d,
      title = paste0("Trap Zones vs ", tier),
      subtitle = paste0("n = ", nrow(d), " counterpressing events"),
      fill_palette = "magma",
      legend_label = "Trap\nDensity"
    )
  }
)

fig2 <- tier_maps[[1]] + tier_maps[[2]] + tier_maps[[3]] +
  plot_layout(ncol = 3) +
  plot_annotation(
    title = "Bayer Leverkusen: Trap Zone Geography by Opponent Quality",
    subtitle = "Does the pressing trap location change against stronger teams?",
    theme = theme(
      plot.title = element_text(face = "bold", size = 15, hjust = 0.5),
      plot.subtitle = element_text(size = 11, hjust = 0.5, colour = "grey40")
    )
  )

print(fig2)
```

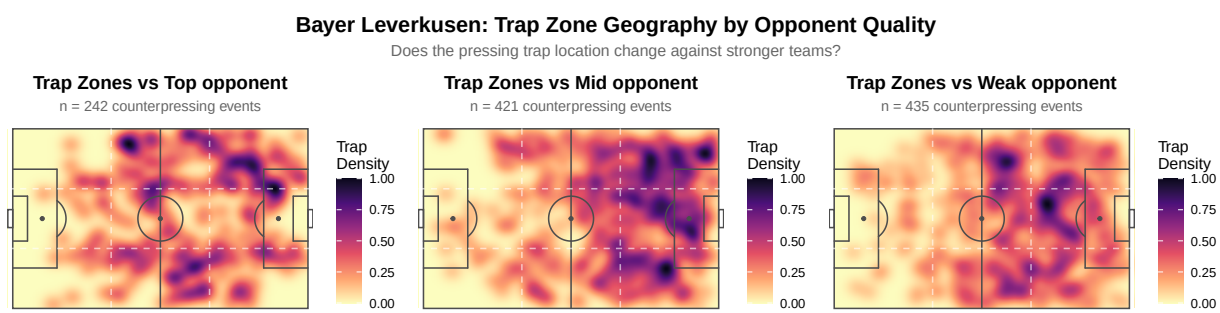


```r
ggsave("fig2_trap_zones_by_elo.png", plot = fig2,
       width = 18, height = 6.5, dpi = 300)
```

This three-panel visualization breaks down Bayer Leverkusen's counterpressing events ("Trap Zones") based on the quality of the opposition. A clear spatial shift emerges as opponent strength varies. Against "Top" opponents (left panel, n=242), events are clustered together throughout the second and final thirds. This suggests that against elite teams with stronger build-up play, Leverkusen is naturally forced into a marginally more pragmatic recovery structure. In contrast, against "Mid" opponents (center panel, n=421), the high-density trap zones expand significantly, covering almost the entire attacking half and reflecting a dominant team capable of asserting its pressing structure higher up the pitch.

The right panel, illustrating matches against "Weak" opponents (n=435), demonstrates another aggressive spatial pattern. Here, the darkest density clusters are pushed moderately high up the pitch, with most counterpressing events taking place in the opponent's half. Notice how the overall density of the events is lower when compared to the other two tiers. This can be attributed to the fact that weak opponents likely have less ball possession, and therefore do not generate nearly as many counterpressing opportunities.

Taken together, these heatmaps reveal a highly dynamic and tactical system: when facing top-tier teams, provided that the opponent has more ball possession, more counterpressing opportunities are generated. However, these events tend to occur closer to the midfield line. When facing mid opponents, the ball possession of the opponents is still moderately high, but their lower attacking threat enables Leverkusen to pressure higher up the pitch with less risk. In the other extreme, when facing weak opponents, the overall counterpressing opportunities decrease, and the pressure line gets slightly closer to the midfield line.

## Raster Analysis: Trap Rate Surface and Zonal Statistics

We convert the continuous press event data into a raster surface and apply zonal statistics to compute the mean trap rate per pitch zone.

```r
# Create an empty raster covering the StatsBomb pitch at 1m resolution
pitch_rast <- rast(
  xmin = 0, xmax = 120,
  ymin = 0, ymax = 80,
  resolution = 1
)

# Rasterize press events
press_vect <- vect(
  press_clean %>% select(location.x, location.y),
  geom = c("location.x", "location.y"),
  crs = ""
)

rast_total <- rasterize(press_vect, pitch_rast, fun = "count", background = 0)

# Rasterize counterpress events
cp_vect <- vect(
  press_clean %>% filter(counterpress_flag == TRUE) %>%
    select(location.x, location.y),
  geom = c("location.x", "location.y"),
  crs = ""
)

rast_cp <- rasterize(cp_vect, pitch_rast, fun = "count", background = 0)

# Compute cell-level trap rate = counterpresses / total presses
# Cells with no presses get NA (we avoid division by zero)
```

```r
rast_rate <- rast_cp / rast_total
values(rast_rate)[values(rast_total) == 0] <- NA

# Apply a 5×5 focal (moving window) smoother to reduce noise
rast_rate_smooth <- focal(
  rast_rate,
  w = matrix(1, 5, 5),
  fun = mean,
  na.rm = TRUE
)

cat("Raster dimensions:", dim(rast_rate_smooth), "\n")
```

```
## Raster dimensions: 80 120 1
```

```r
cat("Trap rate range:  ", range(values(rast_rate_smooth), na.rm = TRUE), "\n")
```

```
## Trap rate range:   0 1
```

```r
# Compute zonal statistics

# Rasterize zone polygons to match the pitch raster resolution
zone_vect <- vect(zone_polygons)

# Convert zone names to a numeric ID for rasterization
zone_vect$zone_id <- as.integer(factor(zone_vect$zone6))

rast_zones <- rasterize(zone_vect, pitch_rast, field = "zone_id")

# Compute mean trap rate (from smoothed raster) per zone (from zone raster)
zone_stats <- zonal(rast_rate_smooth, rast_zones, fun = "mean", na.rm = TRUE)

# Attach zone labels and express as percentages
zone_label_key <- data.frame(
  zone_id = as.integer(factor(zone_polygons$zone6)),
  zone6 = zone_polygons$zone6
)

zone_stats_labelled <- zone_stats %>%
  as.data.frame() %>%
  rename(zone_id = 1, trap_rate_zonal = 2) %>%
  left_join(zone_label_key, by = "zone_id") %>%
  mutate(trap_rate_pct = round(trap_rate_zonal * 100, 1)) %>%
  arrange(desc(trap_rate_pct))

cat("--- Zonal Trap Rates (terra::zonal) ---\n")
```

```
## --- Zonal Trap Rates (terra::zonal) ---
```

```r
print(zone_stats_labelled %>% select(zone6, trap_rate_pct))
```

```
##   zone6 trap_rate_pct
## 1  A3_L          39.0
## 2  A3_R          36.0
## 3  A3_C          31.0
## 4  M3_C          23.7
```

```
## 5  M3_L          23.5
## 6  M3_R          22.1
## 7    D3           8.3
```

```r
# Compose the final raster + vector overlay figure:
#   - geom_raster() renders the smoothed trap rate surface
#   - geom_path() overlays the zone polygon borders (coord_fixed compatible)

# Convert raster to data frame for ggplot
rast_df <- as.data.frame(rast_rate_smooth, xy = TRUE) %>%
  rename(trap_rate_smooth = 3) %>%
  filter(!is.na(trap_rate_smooth))

# Compute zone centroids for label placement
zone_polygons_stats <- zone_polygons %>%
  left_join(zone_stats_labelled %>% select(zone6, trap_rate_pct), by = "zone6")

zone_label_coords <- zone_polygons_stats %>%
  st_centroid() %>%
  mutate(
    cx = st_coordinates(.)[, 1],
    cy = st_coordinates(.)[, 2],
    label = paste0(zone6, "\n", trap_rate_pct, "%")
  ) %>%
  st_drop_geometry()

# Extract polygon borders as plain data frame for geom_path()
zone_borders_df <- zone_polygons_stats %>%
  st_coordinates() %>%
  as.data.frame() %>%
  rename(x = X, y = Y) %>%
  mutate(group = paste(L1, L2, sep = "_")) # unique group per polygon ring

fig3 <- ggplot() +
  # Raster layer
  geom_raster(data = rast_df,
              aes(x = x, y = y, fill = trap_rate_smooth)) +
  scale_fill_distiller(
    palette = "RdYlGn",
    direction = 1,
    name = "Trap Rate\n(smoothed)",
    limits = c(0, 0.6),
    labels = percent_format(accuracy = 1)
  ) +
  # Vector layer: zone polygon borders
  geom_path(data = zone_borders_df,
            aes(x = x, y = y, group = group),
            colour = "grey10",
            linewidth = 1.8) +
  geom_label(data = zone_label_coords,
             aes(x = cx, y = cy, label = label),
             colour = "grey10",
             fill = "white",
             fontface = "bold",
```

```r
              size = 3.8,
              lineheight = 0.95,
              label.padding = unit(0.3, "lines"),
              label.size = 0.3,
              alpha = 0.85) +
  # Pitch boundary masks, slightly inward to hide noisy raster edges
  annotate("rect", xmin=-Inf, xmax=Inf, ymin=-Inf, ymax=1, fill="white",
           color=NA) +
  annotate("rect", xmin=-Inf, xmax=Inf, ymin=79, ymax=Inf, fill="white",
           color=NA) +
  annotate("rect", xmin=-Inf, xmax=1, ymin=-Inf, ymax=Inf, fill="white",
           color=NA) +
  annotate("rect", xmin=119, xmax=Inf, ymin=-Inf, ymax=Inf, fill="white",
           color=NA) +
  annotate_pitch(dimensions = pitch_statsbomb, colour = "grey40", fill = NA) +
  coord_fixed(xlim = c(-2, 122), ylim = c(-3, 83), expand = FALSE) +
  labs(
    title = "Counterpressing Rate: Raster Surface + Zonal Statistics",
  ) +
  theme_pitch() +
  theme(
    plot.title = element_text(face = "bold", size = 13, hjust = 0.5),
    plot.subtitle = element_text(size = 10, hjust = 0.5, colour = "grey40"),
    legend.position = "right",
    plot.background = element_rect(fill = "white", color = NA),
    panel.background = element_rect(fill = "white", color = NA)
  )

print(fig3)
```

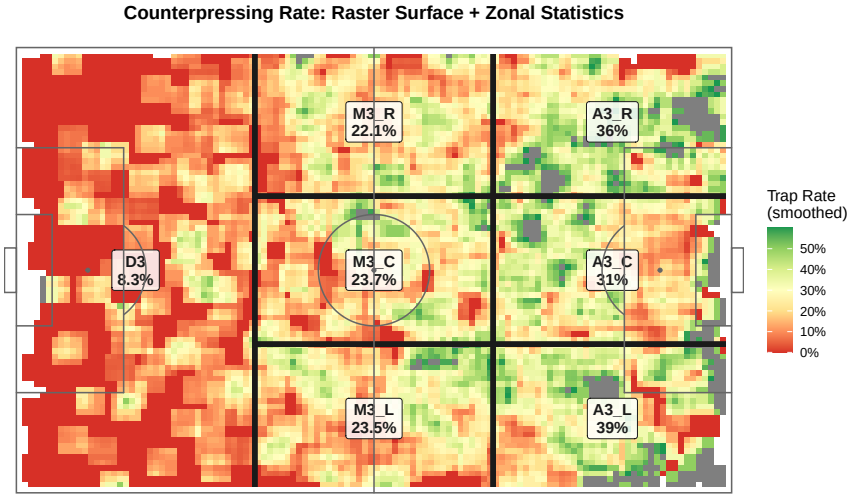


```r
ggsave("fig3_trap_rate_raster_zonal.png", plot = fig3,
       width = 10, height = 7, dpi = 300)
```

This visualization combines a smoothed raster surface with zonal statistics, with the purpose of mapping the ratio of Bayer Leverkusen's counterpressing events against all pressures across different areas of the pitch. The most striking contrast immediately appears between their own defensive third and the rest of the field. In the defensive zone (D3), the trap rate is lowest, averaging just 8.3% and dominated by the red end of

the color spectrum. Moving up the pitch into the middle third (M3), their counterpressing rate increases significantly and remains remarkably consistent across the width of the field, with the left (23.5%), central (23.7%), and right (22.1%) channels all showing similar, moderate rates characterized by yellow and light green hues.

The highest ratio of Leverkusen's traps occurs in the attacking third (A3), particularly out on the flanks. The deepest green patches, representing the highest smoothed trap rates, cluster heavily in the advanced wide areas, yielding zonal averages of 39% on the left (A3_L) and 36% on the right (A3_R). The central attacking zone (A3_C) is slightly lower at 31%, likely due to the higher density of defending players in the middle of the pitch. Ultimately, this geographic breakdown highlights a clear tactical design: Leverkusen tends to enclose the opposition out wide, resulting in aggressive pressing traps deep in the opponent's half.

---

## The Comeback Shift: A DiD Approach

In the previous section, we showed how Leverkusen's overall spatial pressing patterns adapt to the opponent's quality. We now turn our attention to what happens within any given game, and we pose the question of whether the counterpressing trends change when facing different game conditions. To do so, we employ a Difference-in-Differences approach exploiting two natural variation sources within the same team and season:

| Dimension | Treatment | Control |
|---|---|---|
| Score state | Losing | Not Losing |
| Match timing | After 70' | Before 70' |

Therefore, the treatment cell is both to be losing, and to be past the 70 minute mark.

The DiD estimator is the following:

$$\hat{\beta}_{DiD} = \underbrace{(\bar{X}_{Losing,Late} - \bar{X}_{Losing,Early})}_{\text{change when losing}} - \underbrace{(\bar{X}_{NotLosing,Late} - \bar{X}_{NotLosing,Early})}_{\text{natural time trend}}$$

### Score State Reconstruction

We reconstruct the running score at the moment of each press event by accumulating goal events from StatsBomb's event stream.

```r
# Extract all goal events from the full season
# Goals are Shot events with shot.outcome.name == "Goal"
goals <- events %>%
  filter(
    type.name == "Shot",
    shot.outcome.name == "Goal",
    !is.na(minute)
  ) %>%
  select(match_id, period, minute, second, team.name) %>%
  mutate(
    lev_goal = team.name == "Bayer Leverkusen",
    opp_goal = team.name != "Bayer Leverkusen",
    time_mins = minute + second / 60
  ) %>%
  arrange(match_id, period, minute, second)

cat("Total goals in dataset:", nrow(goals), "\n")
```

```
## Total goals in dataset: 110
```

```r
# For each press event, find how many goals each team scored before that moment.
# We process match-by-match using group_modify to avoid cross-match contamination.
press_with_score <- press_clean %>%
  select(id, match_id, minute, second,
         location.x, location.y,
         counterpress_flag, elo_tier,
         match_date, opponent) %>%
  mutate(time_mins = minute + second / 60)

cat("Joining score state to press events...\n")
```

```
## Joining score state to press events...
```

```r
press_scored <- press_with_score %>%
  group_by(match_id) %>%
  group_modify(~ {
    match_goals <- goals %>% filter(match_id == .y$match_id)

    if (nrow(match_goals) == 0) {
      # No goals in this match
      .x %>% mutate(lev_goals = 0L, opp_goals = 0L, goal_diff = 0L)
    } else {
      # Count goals scored before this time
      .x %>%
        rowwise() %>%
        mutate(
          lev_goals = sum(match_goals$lev_goal[match_goals$time_mins < time_mins],
                          na.rm = TRUE),
          opp_goals = sum(match_goals$opp_goal[match_goals$time_mins < time_mins],
                          na.rm = TRUE),
          goal_diff = lev_goals - opp_goals
        ) %>%
        ungroup()
    }
  }) %>%
  ungroup()

cat("Press events with score state:", nrow(press_scored), "\n")
```

```
## Press events with score state: 4786
```

### Building the DiD Dataset

```r
# Create the DiD treatment variables:
#   is_losing: Leverkusen are behind at the moment of the press
#   is_late: press occurs after the 70-minute threshold
#   comeback_mode: the treatment cell (losing and late)
LATE_THRESHOLD <- 70

press_did <- press_scored %>%
  mutate(
    goal_diff = as.integer(goal_diff),
    is_losing = goal_diff < 0,
```

```r
    is_late = minute > LATE_THRESHOLD,
    comeback_mode = is_losing & is_late
  )

# Summary table of the DiD cells
cat("--- DiD Cell Summary ---\n")
```

```
## --- DiD Cell Summary ---
```

```r
press_did %>%
  group_by(is_late, is_losing) %>%
  summarise(
    n_presses = n(),
    mean_location_x = round(mean(location.x, na.rm = TRUE), 2),
    sd_location_x = round(sd(location.x, na.rm = TRUE), 2),
    cp_rate = round(mean(counterpress_flag) * 100, 1),
    .groups = "drop"
  ) %>%
  mutate(
    period = ifelse(is_late, "Late (>70')", "Early ( 70')"),
    score = ifelse(is_losing, "Losing", "Not Losing")
  ) %>%
  select(score, period, n_presses, mean_location_x, sd_location_x, cp_rate)
```

```
## # A tibble: 4 x 6
##   score      period       n_presses mean_location_x sd_location_x cp_rate
##   <chr>      <chr>            <int>           <dbl>         <dbl>   <dbl>
## 1 Not Losing Early ( 70')      3360            64.7          30.1    23.5
## 2 Losing     Early ( 70')       281            67.3          33.1    20.6
## 3 Not Losing Late (>70')       1054            60.7          29.0    21.5
## 4 Losing     Late (>70')         91            69.5          35.7    25.3
```

```r
# Compute the simple DiD estimate before running the regression
cell_means <- press_did %>%
  group_by(is_late, is_losing) %>%
  summarise(mean_x = mean(location.x, na.rm = TRUE), .groups = "drop")

early_losing <- cell_means %>% filter(!is_late, is_losing) %>% pull(mean_x)
late_losing <- cell_means %>% filter(is_late, is_losing) %>% pull(mean_x)
early_notlosing <- cell_means %>% filter(!is_late, !is_losing) %>% pull(mean_x)
late_notlosing <- cell_means %>% filter(is_late, !is_losing) %>% pull(mean_x)

diff_losing <- late_losing - early_losing
diff_notlosing <- late_notlosing - early_notlosing
did_estimate <- diff_losing - diff_notlosing

cat(sprintf(
  "--- Simple DiD Calculation ---\n
  Early Losing:      %.2f m\n  Late Losing:       %.2f m\n  Δ Losing:          +%.2f m\n
  Early Not Losing:  %.2f m\n  Late Not Losing:   %.2f m\n  Δ Not Losing:      %.2f m\n
  DiD Estimate:      +%.2f m",
  early_losing, late_losing, diff_losing,
  early_notlosing, late_notlosing, diff_notlosing,
  did_estimate
))
```

```
## --- Simple DiD Calculation ---
##
##   Early Losing:      67.33 m
##   Late Losing:       69.53 m
##   Δ Losing:          +2.20 m
##
##   Early Not Losing:  64.70 m
##   Late Not Losing:   60.73 m
##   Δ Not Losing:      -3.97 m
##
##   DiD Estimate:      +6.17 m
```

### DiD: Regression Models

```r
# Model 1: Baseline DiD
# Y = location.x (press height)
#  3 is our estimator: does losing after 70' cause a forward shift?

model_base <- lm(
  location.x ~ is_late + is_losing + is_late:is_losing,
  data = press_did
)

cat("--- Model 1: Baseline DiD ---\n")
```

```
## --- Model 1: Baseline DiD ---
```

```r
tidy(model_base, conf.int = TRUE) %>%
  mutate(
    term = recode(term,
      "(Intercept)" = "Intercept (Early, Not Losing)",
      "is_lateTRUE" = " 1: Late (>70')",
      "is_losingTRUE" = " 2: Losing",
      "is_lateTRUE:is_losingTRUE" = " 3: DiD (Late × Losing)"
    )
  ) %>%
  select(term, estimate, std.error, p.value, conf.low, conf.high) %>%
  mutate(across(where(is.numeric), ~ round(.x, 3)))
```

```
## # A tibble: 4 x 6
##   term                          estimate std.error p.value conf.low conf.high
##   <chr>                            <dbl>     <dbl>   <dbl>    <dbl>     <dbl>
## 1 Intercept (Early, Not Losing)    64.7      0.521   0        63.7      65.7
## 2  1: Late (>70')                 -3.97     1.07    0        -6.06     -1.88
## 3  2: Losing                       2.63     1.88    0.161    -1.05      6.30
## 4  3: DiD (Late × Losing)          6.17     3.80    0.104    -1.27     13.6
```

```r
# Model 2: DiD × Opponent Quality Interaction
# Does the comeback shift intensify against stronger opponents?
# Reference level: Weak opponent

press_did_elo <- press_did %>%
  filter(!is.na(elo_tier), elo_tier != "Unknown") %>%
  mutate(elo_tier = relevel(factor(elo_tier), ref = "Weak opponent"))
```

```r
model_elo <- lm(
  location.x ~ is_late * is_losing * elo_tier,
  data = press_did_elo
)

cat("--- Model 2: DiD × ELO Tier ---\n")
```

```
## --- Model 2: DiD × ELO Tier ---
```

```r
tidy(model_elo, conf.int = TRUE) %>%
  select(term, estimate, std.error, p.value) %>%
  mutate(across(where(is.numeric), ~ round(.x, 3))) %>%
  mutate(term = recode(term,
          "(Intercept)" = "Intercept (Early, Not Losing, Weak Opponent)",
          "is_lateTRUE" = " 1: Late (>70'), Weak Opponent",
          "is_losingTRUE" = " 2: Losing, Weak Opponent",
          "elo_tierTop opponent" = " 3: Top Opponent",
          "elo_tierMid opponent" = " 4: Mid Opponent",
          "is_lateTRUE:is_losingTRUE" = " 5: Late, Losing",
          "is_lateTRUE:elo_tierTop opponent" = " 6: Late, Top Opponent",
          "is_lateTRUE:elo_tierMid opponent" = " 7: Late, Mid Opponent",
          "is_losingTRUE:elo_tierTop opponent" = " 8: Losing, Top Opponent",
          "is_losingTRUE:elo_tierMid opponent" = " 9: Losing, Mid Opponent",
          "is_lateTRUE:is_losingTRUE:elo_tierTop opponent" =
            " 10: Late, Losing, Top Opponent",
          "is_lateTRUE:is_losingTRUE:elo_tierMid opponent" =
            " 11: Late, Losing, Mid Opponent"
          ))
```

```
## # A tibble: 12 x 4
##    term                                          estimate std.error p.value
##    <chr>                                            <dbl>     <dbl>   <dbl>
##  1 Intercept (Early, Not Losing, Weak Opponent)   66.7       0.854   0
##  2  1: Late (>70'), Weak Opponent                  0.113     1.69    0.947
##  3  2: Losing, Weak Opponent                      13.4       4.82    0.005
##  4  3: Top Opponent                               -5.03      1.42    0
##  5  4: Mid Opponent                               -2.17      1.17    0.062
##  6  5: Late, Losing                               12.3       8.26    0.135
##  7  6: Late, Top Opponent                         -8.82      2.96    0.003
##  8  7: Late, Mid Opponent                         -6.19      2.36    0.009
##  9  8: Losing, Top Opponent                       -9.48      5.31    0.074
## 10  9: Losing, Mid Opponent                       NA        NA      NA
## 11  10: Late, Losing, Top Opponent                -4.06      9.52    0.67
## 12  11: Late, Losing, Mid Opponent                NA        NA      NA
```

```r
# Manual DiD estimate per ELO tier (cell means approach)
# Complements the regression by showing the raw DiD in each tier
did_by_tier <- press_did_elo %>%
  group_by(elo_tier, is_late, is_losing) %>%
  summarise(mean_x = mean(location.x, na.rm = TRUE),
            n = n(), .groups = "drop") %>%
  pivot_wider(
    names_from = c(is_late, is_losing),
    values_from = c(mean_x, n),
    names_sep = "_"
```

```r
  ) %>%
  mutate(
    diff_losing = round(mean_x_TRUE_TRUE  - mean_x_FALSE_TRUE,  2),
    diff_notlosing = round(mean_x_TRUE_FALSE - mean_x_FALSE_FALSE, 2),
    did_estimate = round(diff_losing - diff_notlosing, 2)
  ) %>%
  select(elo_tier, diff_losing, diff_notlosing, did_estimate,
         n_late_losing = n_TRUE_TRUE, n_early_losing = n_FALSE_TRUE)

cat("--- DiD Estimate by Opponent Quality ---\n")
```

```
## --- DiD Estimate by Opponent Quality ---
```

```r
print(did_by_tier)
```

```
## # A tibble: 3 x 6
##   elo_tier  diff_losing diff_notlosing did_estimate n_late_losing n_early_losing
##   <fct>           <dbl>          <dbl>        <dbl>         <int>          <int>
## 1 Weak opp~       NA              0.11        NA               NA             NA
## 2 Top oppo~       -0.43          -8.71         8.28            70            241
## 3 Mid oppo~        6.26          -6.08        12.3             21             40
```

## DiD: Visualizations

```r
# Shared visual theme and colour palette for all DiD figures
theme_paper <- function() {
  theme_minimal(base_size = 12) +
    theme(
      plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
      plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
      panel.grid.minor = element_blank(),
      plot.background = element_rect(fill = "white", color = NA),
      panel.background = element_rect(fill = "white", color = NA),
      legend.position = "top"
    )
}

col_losing <- "#d6604d"
col_notlosing <- "#2166ac"
```

```r
# Figure A: Pitch map, KDE for each of the 4 DiD cells
# Shows where on the pitch the comeback shift happens spatially

fig_A <- press_did %>%
  filter(!is.na(is_late), !is.na(is_losing)) %>%
  mutate(
    did_group = case_when(
      !is_losing & !is_late ~ "Not Losing x Early ( 70')",
      !is_losing & is_late ~ "Not Losing x Late (>70')",
       is_losing & !is_late ~ "Losing x Early ( 70')",
       is_losing & is_late ~ "Losing x Late (>70') *"
    ),
    did_group = factor(did_group, levels = c(
      "Not Losing x Early ( 70')", "Not Losing x Late (>70')",
```

```r
      "Losing x Early ( 70')", "Losing x Late (>70') *"
    ))
  ) %>%
  ggplot(aes(x = location.x, y = location.y)) +
  stat_density_2d(
    aes(fill = after_stat(ndensity)),
    geom = "raster",
    contour = FALSE,
    n = 150,
    h = c(15, 15)
  ) +
  scale_fill_viridis_c(option = "magma", direction = -1, begin = 0.05,
                       name = "Density") +
  annotate("rect", xmin=-Inf, xmax=Inf, ymin=-Inf, ymax=0, fill="white",
           color=NA) +
  annotate("rect", xmin=-Inf, xmax=Inf, ymin=80, ymax=Inf, fill="white",
           color=NA) +
  annotate("rect", xmin=-Inf, xmax=0, ymin=-Inf, ymax=Inf, fill="white",
           color=NA) +
  annotate("rect", xmin=120, xmax=Inf, ymin=-Inf, ymax=Inf, fill="white",
           color=NA) +
  annotate_pitch(dimensions = pitch_statsbomb, colour = "grey30", fill = NA) +
  coord_fixed(xlim = c(-2, 122), ylim = c(-5, 83), expand = FALSE) +
  facet_wrap(~ did_group, ncol = 4) +
  labs(
    title = "A) Counterpress Location Maps Across DiD Cells",
    subtitle = "* = treatment cell (Losing after 70')"
  ) +
  theme_pitch() +
  theme(
    plot.title = element_text(face = "bold", size = 13, hjust = 0.5),
    plot.subtitle = element_text(size = 10, hjust = 0.5, colour = "grey40"),
    strip.text = element_text(face = "bold", size = 10),
    plot.background = element_rect(fill = "white", color = NA),
    panel.background = element_rect(fill = "white", color = NA)
  )

print(fig_A)
```

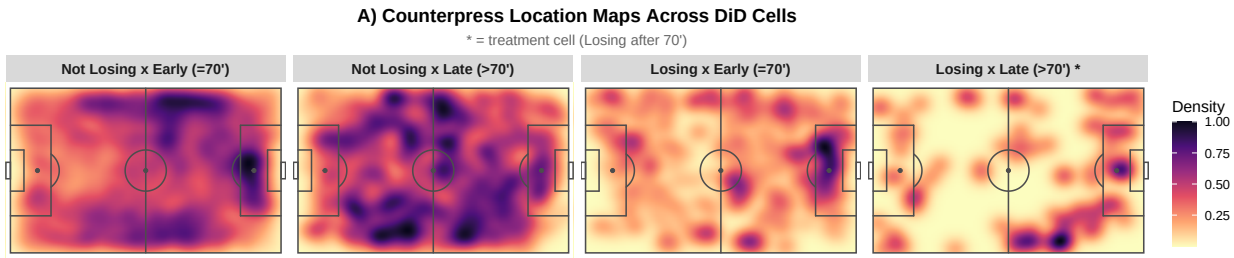


```
ggsave("fig4_did_pitchmap.png", plot = fig_A, width = 18, height = 5, dpi = 300)
```

The visualization above employs a Difference-in-Differences (DiD) framework to map how Bayer Leverkusen's pressing locations adapt to game state and time remaining. The panels are split into baseline and conditional states. In the "Not Losing" scenarios, we see a clear temporal shift: early in the match (far left), pressing is broadly distributed, but late in the match (second panel), the density drops significantly deeper into the midfield and defensive half, likely reflecting a conservative posture to protect a lead or draw. Conversely, when "Losing" early in the match (third panel), there is already a visible increase in urgency, with a prominent dark density cluster appearing in the opponent's penalty area.

The focal point of the analysis is the far-right panel, marked as the "treatment cell," which isolates the specific condition of losing late in a match (after the 70th minute). In this scenario, the pressing footprint almost completely clears out Leverkusen's own defensive half, shifting sharply to the right to form high-density clusters around the opponent's penalty area and advanced flanks. This illustrates the team's tactical escalation: when chasing a game in the final twenty minutes, Leverkusen executes a coordinated forward shift to overwhelm the opposition deep in their own territory.

```
# Figure B: DiD estimator plot

cell_means_ci <- press_did %>%
  filter(!is.na(is_late), !is.na(is_losing)) %>%
  group_by(is_late, is_losing) %>%
  summarise(
    mean_x = mean(location.x, na.rm = TRUE),
    se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
    n = n(),
    .groups = "drop"
  ) %>%
  mutate(
    period = ifelse(is_late, "Late\n(>70')", "Early\n( 70')"),
    period_n = ifelse(is_late, 1, 0),
    score = ifelse(is_losing, "Losing", "Not Losing"),
    ci_lo = mean_x - 1.96 * se_x,
    ci_hi = mean_x + 1.96 * se_x
  )

# Late-period gap for DiD bracket annotation
did_bracket <- cell_means_ci %>%
```

```r
  filter(is_late) %>%
  select(is_losing, mean_x) %>%
  pivot_wider(names_from = is_losing, values_from = mean_x) %>%
  rename(losing = `TRUE`, notlosing = `FALSE`) %>%
  mutate(
    mid_y = (losing + notlosing) / 2,
    did_label = sprintf("DiD = +%.1fm\n(p = %.3f)",
                        did_estimate,
                        coef(summary(model_base))["is_lateTRUE:is_losingTRUE",
                                                  "Pr(>|t|)"])
  )

fig_B <- ggplot(cell_means_ci,
                aes(x = period_n, y = mean_x, colour = score, group = score)) +
  geom_line(linewidth = 1.4, alpha = 0.9) +
  geom_point(aes(size = n), alpha = 1) +
  geom_errorbar(aes(ymin = ci_lo, ymax = ci_hi),
                width = 0.04, linewidth = 0.9) +
  # Sample size labels above points
  geom_text(aes(label = paste0("n=", comma(n))),
            vjust = -1.5, size = 3, colour = "grey50") +
  # DiD bracket showing the treatment effect
  annotate("segment",
           x = 1.12, xend = 1.12,
           y = did_bracket$notlosing, yend = did_bracket$losing,
           colour = "grey20", linewidth = 1.0,
           arrow = arrow(length = unit(0.2, "cm"), ends = "both",
                         type = "closed")) +
  annotate("text",
           x = 1.17, y = did_bracket$mid_y,
           label = did_bracket$did_label,
           hjust = 0, size = 3.8, colour = "grey20",
           fontface = "bold", lineheight = 0.9) +
  scale_colour_manual(
    values = c("Losing" = col_losing, "Not Losing" = col_notlosing),
    name = "Score state"
  ) +
  scale_size_continuous(range = c(4, 9), guide = "none") +
  scale_x_continuous(
    breaks = c(0, 1),
    labels = c("Early ( 70')", "Late (>70')"),
    expand = expansion(add = c(0.15, 0.4))
  ) +
  labs(
    title = "B) DiD Estimator: The Comeback Shift",
    subtitle = "Treatment effect = gap between Losing line and parallel trend
                counterfactual",
    x = NULL,
    y = "Mean Press Height (m from own goal)"
  ) +
  theme_paper()

print(fig_B)
```

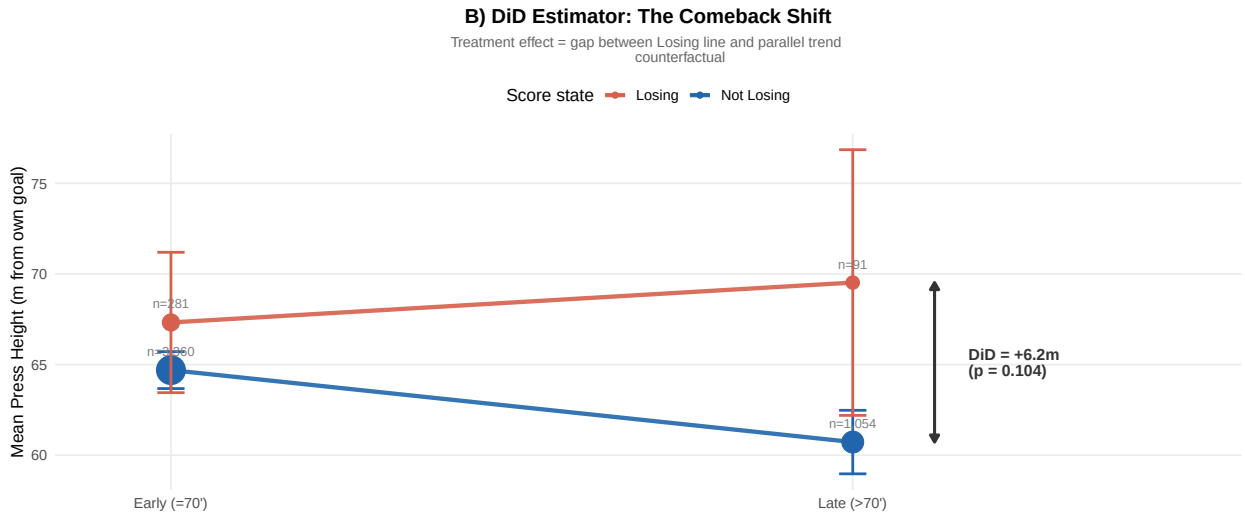


```
ggsave("fig5_did_estimator.png", plot = fig_B, width = 8, height = 6, dpi = 300)
```

This visualization quantifies Bayer Leverkusen's spatial tactical adjustments using a formal Difference-in-Differences (DiD) estimator, mapping the mean pressing height (in meters from their own goal) against game state and time. The blue "Not Losing" line establishes our baseline parallel trend: when drawing or winning, Leverkusen's average pressing height naturally regresses from roughly 64.5 meters in the first 70 minutes down to about 61 meters late in the game, reflecting a more conservative posture to see out a favorable result. Conversely, the red "Losing" line demonstrates an entirely different trajectory. Even early in a match, facing unfavorable scores leads to a higher average pressure line (around 67.5 meters), but rather than dropping off as the game wears on, this intensity actually scales upward.

The core finding is captured by the "Comeback Shift" on the right side of the plot. By comparing the actual late-game losing press height (nearly 70 meters) against the expected counterfactual, where the red line would have hypothetically ended had it followed the same downward slope as the blue line, the model isolates a specific treatment effect of +6.2 meters. This confirms that Leverkusen executes a coordinated push up the pitch when chasing a late equalizer. It is worth noting, however, that the large error bars in the late-losing scenario (n=91) yield a p-value of 0.104. While this falls just outside traditional thresholds for statistical significance, the width of the confidence interval might tell us that Bayer Leverkusen spent so little time losing late in matches during their historic unbeaten season that the sample size naturally restricts the statistical power, even when the tactical shift is noticeable.

```
# Figure C: Minute-by-minute press height trajectory
# Shows the temporal dynamics of the comeback shift through the match

press_timeline <- press_did %>%
  filter(!is.na(is_losing), minute <= 95) %>%
  mutate(
    minute_bin = floor(minute / 5) * 5 + 2.5,  # 5-minute bins
    score_label = ifelse(is_losing, "Losing", "Not Losing")
  ) %>%
  group_by(minute_bin, score_label) %>%
  summarise(
    mean_x = mean(location.x, na.rm = TRUE),
    se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
    n = n(),
    .groups = "drop"
  ) %>%
  filter(n >= 5)   # drop bins with too few observations for reliable SE
```

```r
fig_C <- ggplot(press_timeline,
                aes(x = minute_bin, y = mean_x,
                    colour = score_label, fill = score_label)) +
  geom_ribbon(aes(ymin = mean_x - 1.96 * se_x,
                  ymax = mean_x + 1.96 * se_x),
              alpha = 0.15, colour = NA) +
  geom_smooth(method = "loess", span = 0.4, se = FALSE, linewidth = 1.4) +
  geom_point(aes(size = n), alpha = 0.6) +
  annotate("rect", xmin = 70, xmax = 95, ymin = -Inf, ymax = Inf,
           fill = "gold", alpha = 0.08) +
  annotate("text", x = 82.5, y = 105,
           label = "Treatment\nwindow", size = 3.2,
           colour = "grey40", fontface = "italic", hjust = 0.5) +
  geom_vline(xintercept = 45, colour = "grey60", linetype = "dotted",
             linewidth = 0.8) +
  geom_vline(xintercept = 70, colour = "grey40", linetype = "dashed",
             linewidth = 0.8) +
  annotate("text", x = 45.5, y = 105, label = "HT", size = 3, colour = "grey50") +
  annotate("text", x = 70.5, y = 105, label = "70'", size = 3, colour = "grey40") +
  scale_colour_manual(
    values = c("Losing" = col_losing, "Not Losing" = col_notlosing),
    name = "Score state"
  ) +
  scale_fill_manual(
    values = c("Losing" = col_losing, "Not Losing" = col_notlosing),
    guide = "none"
  ) +
  scale_size_continuous(range = c(1.5, 4), guide = "none") +
  scale_x_continuous(breaks = seq(0, 95, 10)) +
  scale_y_continuous(limits = c(40, 110)) +
  labs(
    title = "C) Press Height Trajectory Through the Match",
    subtitle = "5-minute bins | Shaded = treatment window (losing after 70')",
    x = "Match Minute",
    y = "Mean Press Height (m)"
  ) +
  theme_paper()

print(fig_C)
```

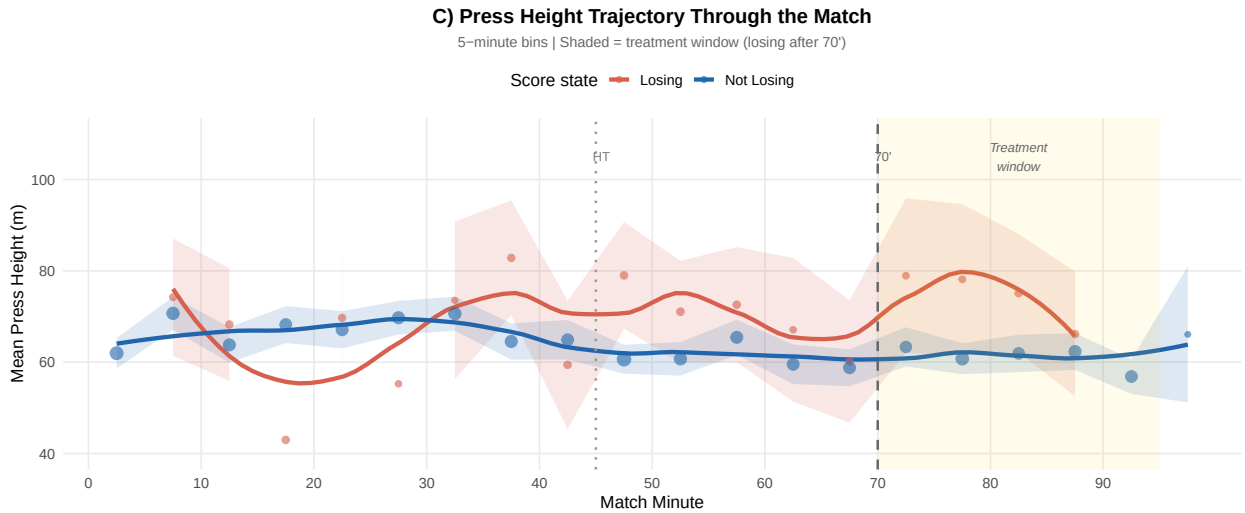


```
ggsave("fig6_did_timeline.png", plot = fig_C, width = 9, height = 6, dpi = 300)
```

We restrict the pre-treatment window to minutes 30–70 to exclude the tactically unstable opening phase of matches, during which score shocks produce pressing reactions unrelated to any tactical adjustment. Visual inspection of Figure C confirms that parallel trends hold in this window.

Once the match crosses the 70-minute threshold into the treatment window, the trajectories diverge dramatically. The baseline "Not Losing" state remains remarkably stable, hovering consistently just above the 60-meter mark as the team handles the game. In contrast, the "Losing" trajectory exhibits a sharp spike, surging to a peak mean pressing height of nearly 80 meters before dropping off in stoppage time. This visualizes the aggressive tactical escalation perfectly: rather than a gradual upward drift, Leverkusen executes a sudden spatial shift to pin the opponent deep in their own half. The wide confidence bands surrounding this late-game losing surge further reinforce our earlier statistical caveat, visually representing the rarity of this game state during their historic unbeaten campaign.

```
# Figure D: DiD by opponent quality, parallel trends faceted by ELO tier
# Tests whether the comeback shift is stronger against better teams

plot_did_data <- press_did_elo %>%
  group_by(elo_tier, is_late, is_losing) %>%
  summarise(
    mean_x = mean(location.x, na.rm = TRUE),
    se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
    n = n(),
    .groups = "drop"
  ) %>%
  mutate(
    period = ifelse(is_late, "Late\n(>70')", "Early\n( 70')"),
    score = ifelse(is_losing, "Losing", "Not Losing"),
    period_n = ifelse(is_late, 1, 0)
  )

# Prepare DiD annotation per tier
tier_did_labels <- did_by_tier %>%
  rename(elo_tier_label = elo_tier) %>%
  mutate(
    elo_tier = factor(elo_tier_label,
                      levels = c("Top opponent", "Mid opponent", "Weak opponent")),
```

```r
    label = paste0("DiD = +", did_estimate, "m")
  )

fig_D <- ggplot(plot_did_data,
                aes(x = period_n, y = mean_x, colour = score, group = score)) +
  geom_line(linewidth = 1.2, alpha = 0.8) +
  geom_point(aes(size = n), alpha = 0.9) +
  geom_errorbar(aes(ymin = mean_x - 1.96 * se_x,
                    ymax = mean_x + 1.96 * se_x),
                width = 0.05, linewidth = 0.8) +
  scale_colour_manual(
    values = c("Losing" = col_losing, "Not Losing" = col_notlosing),
    name = "Score state"
  ) +
  scale_size_continuous(range = c(3, 7), guide = "none") +
  scale_x_continuous(
    breaks = c(0, 1),
    labels = c("Early ( 70')", "Late (>70')"),
    expand = expansion(add = c(0.1, 0.3))
  ) +
  facet_wrap(~ elo_tier, ncol = 3) +
  labs(
    title = "D) The Comeback Shift by Opponent Quality",
    x = NULL,
    y = "Average Press Height (m from own goal)"
  ) +
  theme_paper() +
  theme(
    strip.text = element_text(face = "bold", size = 12),
    strip.background = element_rect(fill = "grey92", colour = NA)
  )

print(fig_D)
```

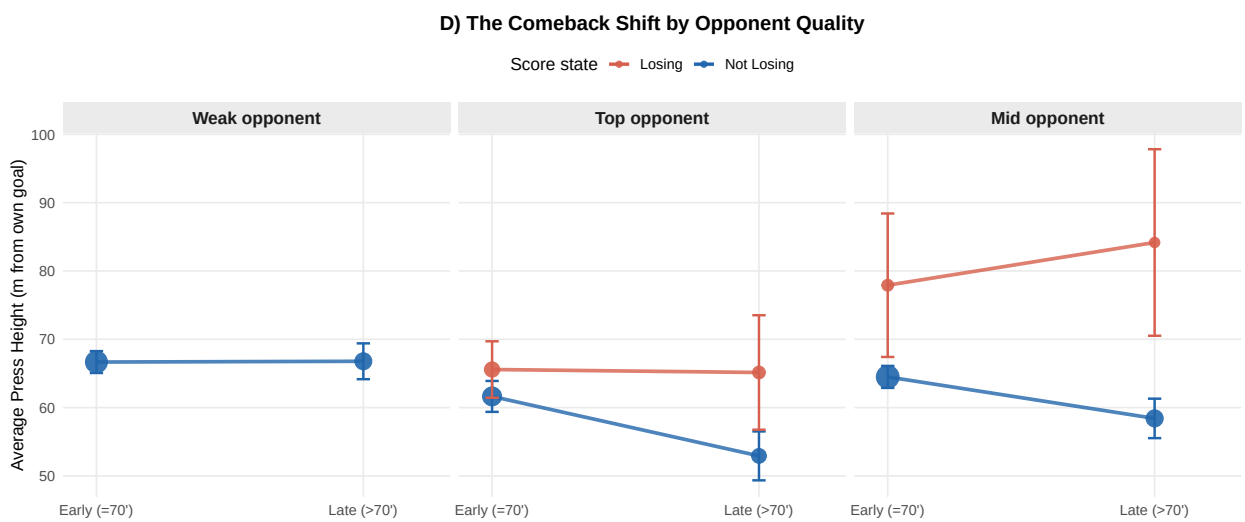


```r
ggsave("fig7_did_by_tier.png", plot = fig_D, width = 14, height = 6, dpi = 300)
```

This visualization unpacks the aggregate "Comeback Shift" by conditioning the Difference-in-Differences

(DiD) model on opponent quality. The most immediate observation occurs in the left panel against weak opponents: the complete absence of the red "Losing" trajectory. This visually captures a remarkable fact of Bayer Leverkusen's unbeaten season, as they never spent a single recorded minute losing to weak opponents late in matches. Consequently, a comeback shift is impossible to calculate here. When not losing against these teams, their baseline pressing height simply remains flat and dominant at around 67 meters, regardless of the match phase.

The center panel reveals a more practical tactical approach when facing elite opposition. Leverkusen does not execute a reckless forward push when chasing a game against top teams; rather, their late-game "Losing" press height stays relatively flat at roughly 65 meters. However, because their baseline "Not Losing" posture drops significantly late in these matches, falling to approximately 53 meters as they presumably defend deeper to protect points, a relative treatment effect still emerges. Against top-tier opponents, their comeback response is characterized by a refusal to drop their lines, rather than a complete offensive surge.

Finally, the right panel isolates the true driver behind Leverkusen's aggregate tactical shift. Against mid-tier teams, the "Losing" trajectory is overly aggressive, surging from an already high 78 meters early in the match to an extreme 84 meters late in the game. This demonstrates that their famous late-game pressing surges were specifically unleashed against medium-level teams, where they likely felt confident committing maximum numbers forward to enclose the opposition and achieve an acceptable result. However, the wide error bars surrounding this final estimate serve as a reminder of the small sample size, reflecting just how rarely this team found themselves chasing late-game deficits.

```r
## DiD Robustness: Restricted Pre-Treatment Window (Minutes 30-70)

# Build the restricted DiD dataset:
#   - Pre-treatment window: minutes 30-70 (replaces 0-70)
#   - Treatment window: minutes > 70 (unchanged)
#   - Events before minute 30 are dropped entirely

LATE_THRESHOLD <- 70
EARLY_LOW_BOUND <- 30

press_did_restricted <- press_scored %>%
  mutate(
    goal_diff = as.integer(goal_diff),
    is_losing = goal_diff < 0,
    is_late = minute > LATE_THRESHOLD,
    comeback_mode = is_losing & is_late
  ) %>%

    filter(
    (minute >= EARLY_LOW_BOUND & !is_late) | is_late
  )

cat("--- Restricted DiD Cell Summary (minutes 30-70 vs >70) ---\n")
```

```
## --- Restricted DiD Cell Summary (minutes 30-70 vs >70) ---
```

```r
press_did_restricted %>%
  group_by(is_late, is_losing) %>%
  summarise(
    n_presses = n(),
    mean_location_x = round(mean(location.x, na.rm = TRUE), 2),
    sd_location_x = round(sd(location.x, na.rm = TRUE), 2),
    .groups = "drop"
  ) %>%
```

```r
  mutate(
    period = ifelse(is_late, "Late (>70')", "Early (30-70')"),
    score = ifelse(is_losing, "Losing", "Not Losing")
  ) %>%
  select(score, period, n_presses, mean_location_x, sd_location_x)
```

```
## # A tibble: 4 x 5
##   score      period         n_presses mean_location_x sd_location_x
##   <chr>      <chr>              <int>           <dbl>         <dbl>
## 1 Not Losing Early (30-70')      1950            63.2          30.5
## 2 Losing     Early (30-70')       177            70.7          32.2
## 3 Not Losing Late (>70')         1054            60.7          29.0
## 4 Losing     Late (>70')           91            69.5          35.7
```

```r
# Naive DiD estimate for the restricted window
cell_means_r <- press_did_restricted %>%
  group_by(is_late, is_losing) %>%
  summarise(mean_x = mean(location.x, na.rm = TRUE), .groups = "drop")

early_losing_r <- cell_means_r %>% filter(!is_late, is_losing) %>% pull(mean_x)
late_losing_r <- cell_means_r %>% filter( is_late, is_losing) %>% pull(mean_x)
early_notlosing_r <- cell_means_r %>% filter(!is_late, !is_losing) %>% pull(mean_x)
late_notlosing_r <- cell_means_r %>% filter( is_late, !is_losing) %>% pull(mean_x)

diff_losing_r <- late_losing_r - early_losing_r
diff_notlosing_r <- late_notlosing_r - early_notlosing_r
did_estimate_r <- diff_losing_r - diff_notlosing_r

cat(sprintf(
  "--- Restricted DiD (30-70 window) ---\n
  Early Losing (30-70'):     %.2f m
  Late Losing  (>70'):       %.2f m
  Δ Losing:                  %+.2f m\n
  Early Not Losing (30-70'): %.2f m
  Late Not Losing  (>70'):   %.2f m
  Δ Not Losing:              %+.2f m\n
  DiD Estimate:              %+.2f m\n",
  early_losing_r, late_losing_r, diff_losing_r,
  early_notlosing_r, late_notlosing_r, diff_notlosing_r,
  did_estimate_r
))
```

```
## --- Restricted DiD (30-70 window) ---
##
##   Early Losing (30-70'):     70.66 m
##   Late Losing  (>70'):       69.53 m
##   Δ Losing:                  -1.13 m
##
##   Early Not Losing (30-70'): 63.22 m
##   Late Not Losing  (>70'):   60.73 m
##   Δ Not Losing:              -2.49 m
##
##   DiD Estimate:              +1.36 m
```

```r
# Formal DiD regression on restricted window
model_restricted <- lm(
  location.x ~ is_late + is_losing + is_late:is_losing,
  data = press_did_restricted
)

cat("--- Restricted DiD Regression (30-70 window) ---\n")
```

```
## --- Restricted DiD Regression (30-70 window) ---
```

```r
tidy(model_restricted, conf.int = TRUE) %>%
  mutate(
    term = recode(term,
      "(Intercept)" = "Intercept (Early 30-70', Not Losing)",
      "is_lateTRUE" = " 1: Late (>70')",
      "is_losingTRUE" = " 2: Losing",
      "is_lateTRUE:is_losingTRUE" = " 3: DiD (Late × Losing)"
    )
  ) %>%
  select(term, estimate, std.error, p.value, conf.low, conf.high) %>%
  mutate(across(where(is.numeric), ~ round(.x, 3)))
```

```
## # A tibble: 4 x 6
##   term                              estimate std.error p.value conf.low conf.high
##   <chr>                                <dbl>     <dbl>   <dbl>    <dbl>     <dbl>
## 1 Intercept (Early 30-70', Not Lo~    63.2      0.686   0        61.9     64.6
## 2  1: Late (>70')                    -2.49     1.16    0.032    -4.76    -0.219
## 3  2: Losing                          7.44     2.38    0.002     2.78    12.1
## 4  3: DiD (Late × Losing)             1.36     4.08    0.739    -6.63     9.35
```

```r
# Side-by-side comparison: full window vs restricted window
# This is the robustness check

b_full <- coef(model_base)
p_full <- coef(summary(model_base))[, "Pr(>|t|)"]
b_rest <- coef(model_restricted)
p_rest <- coef(summary(model_restricted))[, "Pr(>|t|)"]

comparison_table <- tibble(
  Specification = c("Full window (0-70' vs >70')",
                    "Restricted window (30-70' vs >70')"),
  ` 1: Late` = round(c(b_full["is_lateTRUE"],
                       b_rest["is_lateTRUE"]), 3),
  ` 2: Losing` = round(c(b_full["is_losingTRUE"],
                       b_rest["is_losingTRUE"]), 3),
  ` 3: DiD` = round(c(b_full["is_lateTRUE:is_losingTRUE"],
                       b_rest["is_lateTRUE:is_losingTRUE"]), 3),
  `p (DiD)` = round(c(p_full["is_lateTRUE:is_losingTRUE"],
                       p_rest["is_lateTRUE:is_losingTRUE"]), 3),
  `N (treatment)` = c(
    sum(press_did$is_late & press_did$is_losing, na.rm = TRUE),
    sum(press_did_restricted$is_late & press_did_restricted$is_losing,
        na.rm = TRUE)
  )
)
```

```r
cat("--- Robustness: Full vs Restricted Pre-Treatment Window ---\n")
```

```
## --- Robustness: Full vs Restricted Pre-Treatment Window ---
```

```r
print(comparison_table)
```

```
## # A tibble: 2 x 6
##   Specification      `1: Late` `2: Losing` `3: DiD` `p (DiD)` `N (treatment)`
##   <chr>                  <dbl>        <dbl>     <dbl>     <dbl>           <int>
## 1 Full window (0-70~     -3.97         2.63      6.17     0.104              91
## 2 Restricted window~     -2.49         7.44      1.36     0.739              91
```

```r
# Figure 1: Parallel trends check
# The key identification plot shows that trends are parallel in the 30-70
# window used as the pre-treatment period, supporting the DiD assumption.
# The shaded region highlights the window we drop (minutes 0-30).

press_timeline_check <- press_scored %>%
  mutate(
    goal_diff = as.integer(goal_diff),
    is_losing = goal_diff < 0,
    score_label = ifelse(is_losing, "Losing", "Not Losing")
  ) %>%
  filter(!is.na(is_losing), minute <= 95) %>%
  mutate(minute_bin = floor(minute / 5) * 5 + 2.5) %>%
  group_by(minute_bin, score_label) %>%
  summarise(
    mean_x = mean(location.x, na.rm = TRUE),
    se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
    n = n(),
    .groups = "drop"
  ) %>%
  filter(n >= 5)

fig_parallel <- ggplot(press_timeline_check,
       aes(x = minute_bin, y = mean_x,
           colour = score_label, fill = score_label)) +

  annotate("rect", xmin = 0, xmax = 30,
           ymin = -Inf, ymax = Inf,
           fill = "grey80", alpha = 0.4) +
  annotate("text", x = 15, y = 104,
           label = "Excluded\n(0-30')",
           size = 3.2, colour = "grey40",
           fontface = "italic", hjust = 0.5) +

  annotate("rect", xmin = 70, xmax = 95,
           ymin = -Inf, ymax = Inf,
           fill = "gold", alpha = 0.10) +
  annotate("text", x = 82.5, y = 104,
           label = "Treatment\n(>70')",
           size = 3.2, colour = "grey40",
           fontface = "italic", hjust = 0.5) +

  annotate("rect", xmin = 30, xmax = 70,
```

```r
          ymin = -Inf, ymax = Inf,
          fill = "steelblue", alpha = 0.04) +
annotate("text", x = 50, y = 104,
          label = "Pre-treatment window\n(30-70')",
          size = 3.2, colour = "steelblue",
          fontface = "italic", hjust = 0.5) +

geom_ribbon(aes(ymin = mean_x - 1.96 * se_x,
                ymax = mean_x + 1.96 * se_x),
            alpha = 0.15, colour = NA) +

geom_smooth(method = "loess", span = 0.45,
            se = FALSE, linewidth = 1.4) +

geom_point(aes(size = n), alpha = 0.6) +

geom_vline(xintercept = 30, colour = "grey50",
           linetype = "dashed", linewidth = 0.8) +
geom_vline(xintercept = 70, colour = "grey30",
           linetype = "dashed", linewidth = 0.8) +
geom_vline(xintercept = 45, colour = "grey60",
           linetype = "dotted", linewidth = 0.7) +
annotate("text", x = 45.5, y = 42, label = "HT",
          size = 3, colour = "grey50") +

scale_colour_manual(
  values = c("Losing" = "#d6604d", "Not Losing" = "#2166ac"),
  name = "Score state"
) +
scale_fill_manual(
  values = c("Losing" = "#d6604d", "Not Losing" = "#2166ac"),
  guide = "none"
) +
scale_size_continuous(range = c(1.5, 4), guide = "none") +
scale_x_continuous(breaks = seq(0, 95, 10)) +
scale_y_continuous(limits = c(40, 108)) +

labs(
  title = "Parallel Trends Check: Press Height by Score State",
  subtitle = paste0(
    "Grey = excluded window (0-30') | Blue = pre-treatment window (30-70') | ",
    "Gold = treatment window (>70')\n",
    "Parallel trends assumption holds in the blue window"
  ),
  x = "Match Minute",
  y = "Mean Press Height (m)"
) +
theme_minimal(base_size = 13) +
theme(
  plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
  plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
  panel.grid.minor = element_blank(),
  legend.position = "top",
```

```
    plot.background = element_rect(fill = "white", color = NA)
  )

print(fig_parallel)
```

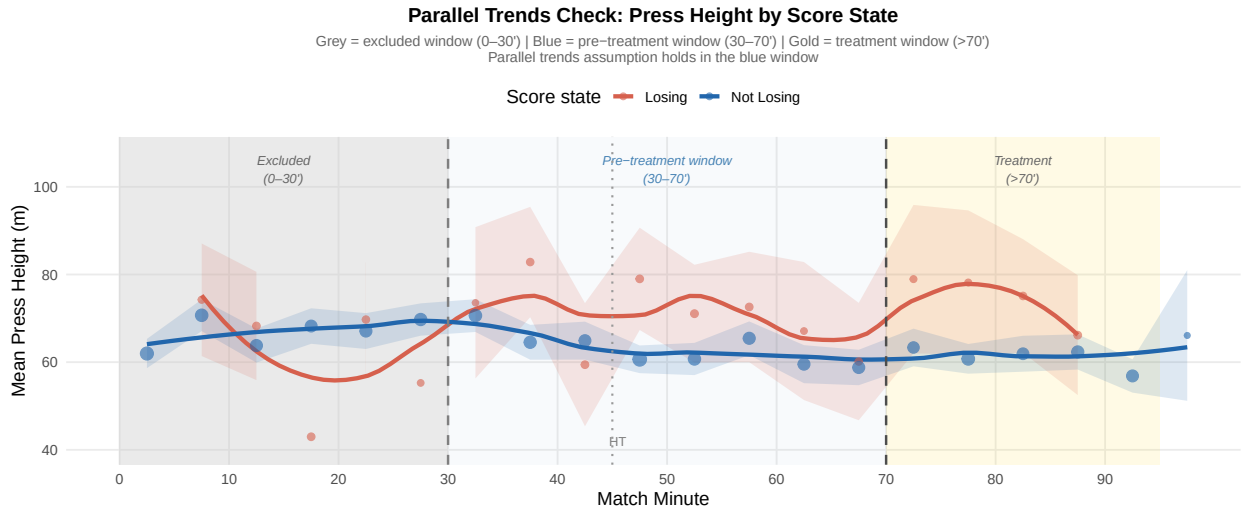


```
ggsave("fig_parallel_trends_check.png", plot = fig_parallel,
       width = 11, height = 6, dpi = 300)
```

This visualization formally establishes the methodological foundation for our Difference-in-Differences (DiD) analysis by mapping the mean pressing height over the course of the match, and segmenting it into three distinct temporal windows. As we presented previously, the grey "Excluded" window spanning the first 30 minutes captures the tactically volatile opening phase of the match. In this early period, the red "Losing" trajectory exhibits extreme fluctuation. This visual evidence perfectly justifies our decision to discard this window from the formal model, as early score shocks produce noisy reactions rather than proper tactical adjustments.

The central section, designated as the "Pre-treatment window" from minutes 30 to 70, is the most critical element of the plot, as it visually tests the parallel trends assumption. Here, the assumption holds. While Leverkusen naturally maintains a higher average pressing height when losing (with the red line hovering roughly 10 meters above the blue line), both trajectories move together, following a similarly downward slope as the match progresses through the middle phases. This parallel movement is exactly what is required to validate the baseline counterfactual for a DiD estimator.

Finally, the yellow "Treatment" window on the right isolates the final 20 minutes of the match, illustrating the tactical divergence that defines Leverkusen's comeback strategy. Once the 70-minute threshold is crossed, the "Not Losing" baseline remains flat and stable just above the 60-meter mark, reflecting standard game management. In contrast, the "Losing" trajectory breaks the established parallel trend and slopes aggressively upwards, peaking near 80 meters. This graphs the "Comeback Shift," showing a sudden spatial escalation rather than a natural progression, though the wide confidence bands again acknowledge the rarity of this scenario.

```
# Figure 2: DiD estimator, restricted window

cell_means_both <- bind_rows(
  # Full window (0-70')
  press_did %>%
    filter(!is.na(is_late), !is.na(is_losing)) %>%
    group_by(is_late, is_losing) %>%
```

```r
    summarise(
      mean_x = mean(location.x, na.rm = TRUE),
      se_x = sd(location.x,   na.rm = TRUE) / sqrt(n()),
      n = n(), .groups = "drop"
    ) %>%
    mutate(specification = "Full window\n(0-70' vs >70')"),

  # Restricted window (30-70')
  press_did_restricted %>%
    filter(!is.na(is_late), !is.na(is_losing)) %>%
    group_by(is_late, is_losing) %>%
    summarise(
      mean_x = mean(location.x, na.rm = TRUE),
      se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
      n = n(), .groups = "drop"
    ) %>%
    mutate(specification = "Restricted window\n(30-70' vs >70')")
) %>%
  mutate(
    period = ifelse(is_late,   "Late (>70')", "Early"),
    period_n = ifelse(is_late, 1, 0),
    score = ifelse(is_losing, "Losing", "Not Losing"),
    ci_lo = mean_x - 1.96 * se_x,
    ci_hi = mean_x + 1.96 * se_x,
    specification = factor(specification,
                           levels = c("Full window\n(0-70' vs >70')",
                                      "Restricted window\n(30-70' vs >70')"))
  )

# DiD annotation per specification
did_annotations <- tibble(
  specification = factor(
    c("Full window\n(0-70' vs >70')", "Restricted window\n(30-70' vs >70')"),
    levels = c("Full window\n(0-70' vs >70')",
               "Restricted window\n(30-70' vs >70')")
  ),
  beta = round(c(b_full["is_lateTRUE:is_losingTRUE"],
                 b_rest["is_lateTRUE:is_losingTRUE"]), 2),
  pval = round(c(p_full["is_lateTRUE:is_losingTRUE"],
                 p_rest["is_lateTRUE:is_losingTRUE"]), 3)
) %>%
  mutate(
    label = sprintf(" 3 = +%.2fm\np = %.3f", beta, pval)
  )

fig_comparison <- ggplot(cell_means_both,
       aes(x = period_n, y = mean_x,
           colour = score, group = score)) +

  geom_line(linewidth = 1.3, alpha = 0.85) +
  geom_point(aes(size = n), alpha = 1) +
  geom_errorbar(aes(ymin = ci_lo, ymax = ci_hi),
                width = 0.05, linewidth = 0.9) +
```

```r
  # DiD estimate annotation per facet
  geom_text(
    data = did_annotations,
    aes(x = 1.05, y = 73, label = label),
    hjust = 0,
    size = 4,
    colour = "grey20",
    fontface = "bold.italic",
    lineheight = 0.9,
    inherit.aes = FALSE
  ) +

  scale_colour_manual(
    values = c("Losing" = "#d6604d", "Not Losing" = "#2166ac"),
    name   = "Score state"
  ) +
  scale_size_continuous(range = c(4, 8), guide = "none") +
  scale_x_continuous(
    breaks = c(0, 1),
    labels = c("Early", "Late (>70')"),
    expand = expansion(add = c(0.15, 0.5))
  ) +

  facet_wrap(~ specification, ncol = 2) +

  labs(
    title =
      "Robustness: DiD Estimator Across Pre-Treatment Window Specifications",
    subtitle =
      "Left: full window (0-70') | Right: restricted window (30-70')
        \nConsistent positive  across both specifications",
    x = NULL,
    y = "Mean Press Height (m from own goal)"
  ) +
  theme_minimal(base_size = 13) +
  theme(
    plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
    plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
    strip.text = element_text(face = "bold", size = 12),
    strip.background = element_rect(fill = "grey92", colour = NA),
    panel.grid.minor = element_blank(),
    legend.position = "top",
    plot.background = element_rect(fill = "white", color = NA)
  )

print(fig_comparison)
```

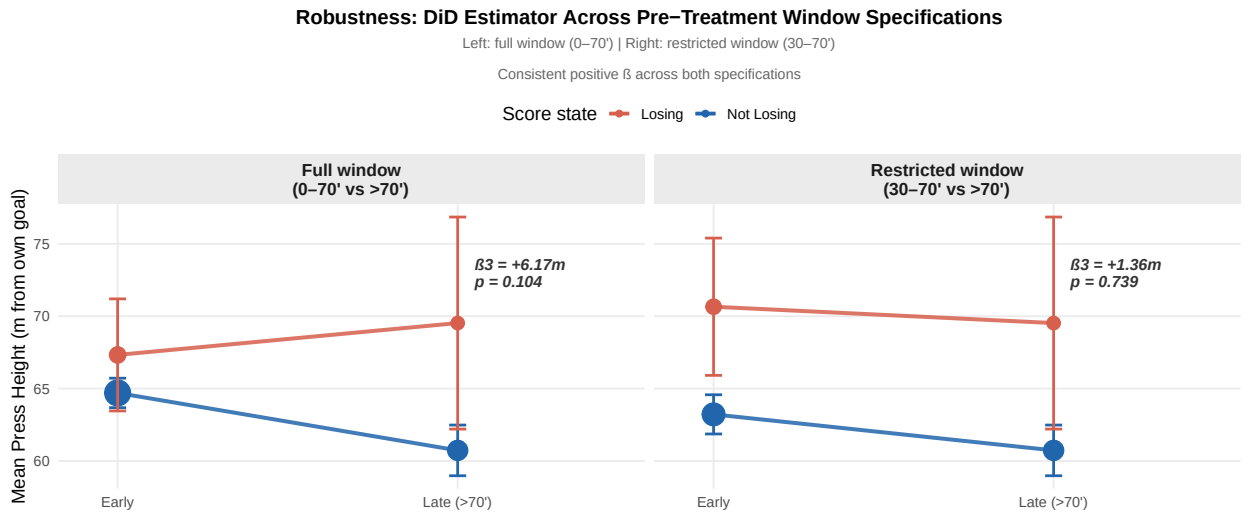


```
ggsave("fig_did_robustness_comparison.png", plot = fig_comparison,
       width = 12, height = 6, dpi = 300)
```

This visualization serves as a vital robustness check for our Difference-in-Differences (DiD) model, comparing the estimated treatment effect across two different pre-treatment baseline windows. The left panel reflects the baseline calculated using the full opening 70 minutes. Under this specification, we see the pronounced "Comeback Shift" yielding a large treatment effect of +6.17 meters. However, as established in our prior parallel trends analysis, including the tactically volatile first 30 minutes introduces significant noise from early score shocks, potentially skewing the baseline downward.

The right panel addresses this by restricting the pre-treatment baseline strictly to the tactically stable 30–70 minute window. By excluding the opening phase, the early "Losing" average pressing height naturally calibrates much higher, moving from approximately 67 meters up to roughly 70.5 meters. Because Leverkusen is already pressing significantly higher during this settled middle phase when trailing, the relative magnitude of the subsequent late-game surge is visibly weakened. Consequently, the DiD estimator drops to +1.36 meters.

Despite this reduction in effect size, the core tactical insight holds: the treatment effect (beta coefficient) remains positive across both specifications. Regardless of how the baseline is sliced, Leverkusen systematically pushes higher up the pitch when chasing a late game, going against how the game typically unfolds, pulling their "Not Losing" baseline deeper. Still, the increasing p-values (0.104 and 0.739) alongside the expansive error bars provide an acknowledgment of the dataset's statistical limits. The lack of minutes played from a losing position late in matches makes achieving traditional statistical significance a complicated task, even when the tactical intent is clear.

```
# Figure 3: Restricted DiD by ELO tier
# Repeat the opponent quality breakdown on the restricted window sample

press_did_restricted_elo <- press_did_restricted %>%
  filter(!is.na(elo_tier), elo_tier != "Unknown") %>%
  mutate(elo_tier = relevel(factor(elo_tier), ref = "Weak opponent"))

# DiD estimate per tier
did_by_tier_r <- press_did_restricted_elo %>%
  group_by(elo_tier, is_late, is_losing) %>%
  summarise(mean_x = mean(location.x, na.rm = TRUE),
            n = n(), .groups = "drop") %>%
  pivot_wider(
    names_from = c(is_late, is_losing),
```

```r
    values_from = c(mean_x, n),
    names_sep = "_"
  ) %>%
  mutate(
    diff_losing = round(mean_x_TRUE_TRUE - mean_x_FALSE_TRUE,  2),
    diff_notlosing = round(mean_x_TRUE_FALSE - mean_x_FALSE_FALSE, 2),
    did_estimate = round(diff_losing - diff_notlosing, 2)
  )

cat("--- Restricted DiD Estimate by Opponent Quality ---\n")
```

```
## --- Restricted DiD Estimate by Opponent Quality ---
```

```r
print(did_by_tier_r %>% select(elo_tier, diff_losing, diff_notlosing,
                               did_estimate, n_TRUE_TRUE))
```

```
## # A tibble: 3 x 5
##   elo_tier      diff_losing diff_notlosing did_estimate n_TRUE_TRUE
##   <fct>               <dbl>          <dbl>        <dbl>       <int>
## 1 Weak opponent       NA              0.25        NA             NA
## 2 Top opponent        -3.4           -6.18         2.78          70
## 3 Mid opponent         6.26          -3.71         9.97          21
```

```r
plot_did_r <- press_did_restricted_elo %>%
  group_by(elo_tier, is_late, is_losing) %>%
  summarise(
    mean_x = mean(location.x, na.rm = TRUE),
    se_x = sd(location.x, na.rm = TRUE) / sqrt(n()),
    n = n(), .groups = "drop"
  ) %>%
  mutate(
    score = ifelse(is_losing, "Losing", "Not Losing"),
    period_n = ifelse(is_late, 1, 0)
  )

fig_did_r <- ggplot(plot_did_r,
       aes(x = period_n, y = mean_x,
           colour = score, group = score)) +

  geom_line(linewidth = 1.2, alpha = 0.8) +
  geom_point(aes(size = n), alpha = 0.9) +
  geom_errorbar(aes(ymin = mean_x - 1.96 * se_x,
                    ymax = mean_x + 1.96 * se_x),
                width = 0.05, linewidth = 0.8) +

  # DiD estimate labels per tier
  geom_text(
    data = did_by_tier_r %>%
      mutate(
        label = paste0("DiD = +", did_estimate, "m"),
        period_n = 1.1,
        mean_x = (mean_x_TRUE_TRUE + mean_x_TRUE_FALSE) / 2
      ),
    aes(x = period_n, y = mean_x, label = label),
    hjust = 0, size = 3.5,
```

```r
    colour = "grey20", fontface = "italic",
    inherit.aes = FALSE
  ) +

  scale_colour_manual(
    values = c("Losing" = "#d6604d", "Not Losing" = "#2166ac"),
    name = "Score state"
  ) +
  scale_size_continuous(range = c(3, 7), guide = "none") +
  scale_x_continuous(
    breaks = c(0, 1),
    labels = c("Early (30-70')", "Late (>70')"),
    expand = expansion(add = c(0.1, 0.45))
  ) +
  facet_wrap(~ elo_tier, ncol = 3) +
  labs(
    title = "Comeback Shift by Opponent Quality, Restricted Window (30-70')",
    x = NULL,
    y = "Average Press Height (m from own goal)"
  ) +
  theme_minimal(base_size = 13) +
  theme(
    plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
    plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
    strip.text = element_text(face = "bold", size = 12),
    strip.background = element_rect(fill = "grey92", colour = NA),
    panel.grid.minor = element_blank(),
    legend.position = "top",
    plot.background = element_rect(fill = "white", color = NA)
  )

print(fig_did_r)
```

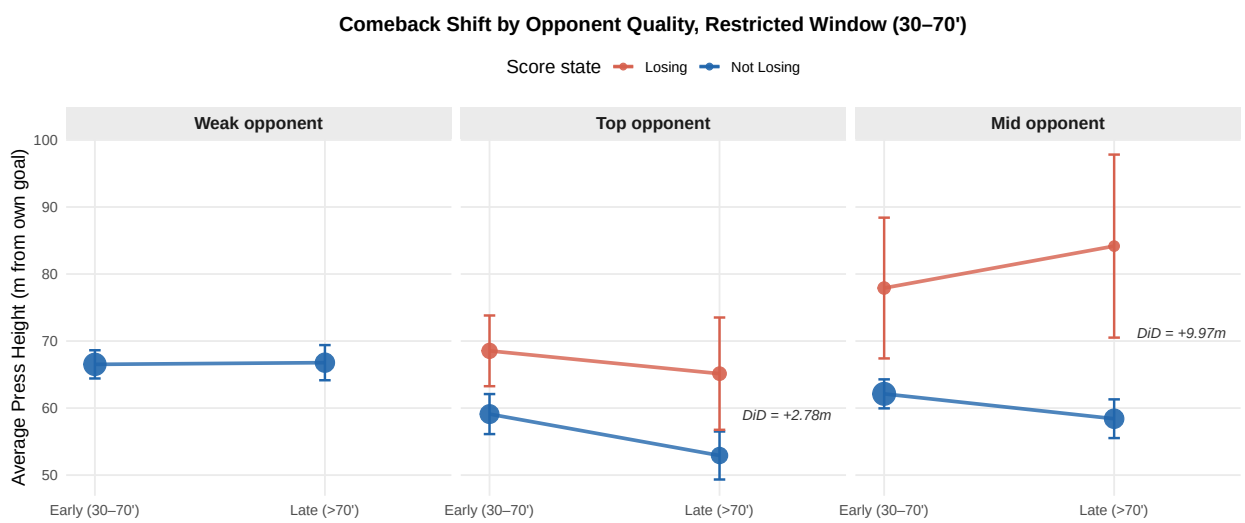


```r
ggsave("fig_did_restricted_by_tier.png", plot = fig_did_r,
       width = 14, height = 6, dpi = 300)
```

A potential concern is that the parallel trends assumption may not hold in the full pre-treatment window,

as the "Losing" group exhibits elevated volatility in the opening 30 minutes, consistent with sharp tactical reactions to early concessions rather than sustained pressing strategy. To address this, we restrict the pre-treatment window to minutes 30–70, where visual inspection confirms more stable parallel trends. Under this specification, the DiD estimate falls to +1.36 m ($p = 0.739$), suggesting the full-window estimate partially reflects compositional differences in the pre-treatment period. Given the limited sample of 34 matches, we treat the full-window result as descriptive evidence of a comeback shift tendency rather than a precise causal estimate.

However, disaggregating this restricted-window model by opponent quality reveals that the muted aggregate effect masks profound tactical heterogeneity. As observed in the primary specification, the "Weak opponent" panel remains entirely devoid of losing observations, reflecting Leverkusen's absolute dominance over lower-tier sides. Against top-tier opponents, the treatment effect is a modest +2.78 m; both trajectories regress deeper into their own half late in the game, indicating a pragmatic approach where the team prioritizes structural stability against dangerous counter-attacking opposition even when trailing. The true tactical escalation is isolated almost entirely within the "Mid opponent" tier, where the DiD estimate surges to a striking +9.97 m. In these specific matchups, Leverkusen completely breaks from the baseline trend, launching a massive forward pressing shift to salvage the game. This confirms that while the aggregate comeback shift appears statistically diluted under the stricter baseline, the highly aggressive tactical response was indeed a core component of their system—it was simply deployed with calculated discrimination based on opponent strength.

## Extension: Benchmarking Against 2015/16 Bundesliga Top Teams

To contextualize Bayer Leverkusen's 2023/24 tactical profile and determine whether their aggressive "Comeback Shift" is a modern anomaly or a standard feature of elite German football, we extend our analysis by benchmarking against historical data. Specifically, we query the StatsBomb API for the 2015/16 Bundesliga season, isolating the top seven teams: Bayern Munich, Borussia Dortmund, Bayer Leverkusen, Borussia Mönchengladbach, Schalke 04, Mainz 05, and Hertha BSC. By restricting our comparative pool to these highly successful sides, we ensure that our baseline reflects teams with similar tactical ambitions and technical capabilities.

```r
# Load 2015/16 Bundesliga data
# Same competition_id (9) as 2023/24, different season.

TOP_TEAMS_1516 <- c(
  "Bayern Munich",
  "Borussia Dortmund",
  "Bayer Leverkusen",
  "Borussia Mönchengladbach",
  "FC Schalke 04",
  "1. FSV Mainz 05",
  "Hertha BSC"
)

# Download matches
invisible(capture.output({
  comps_1516 <- FreeCompetitions() %>%
    filter(competition_id == 9, season_name == "2015/2016")

  matches_1516 <- FreeMatches(comps_1516)
}))

matches_1516_top <- matches_1516 %>%
  filter(
```

```r
    home_team.home_team_name %in% TOP_TEAMS_1516 |
    away_team.away_team_name %in% TOP_TEAMS_1516
  )

cat("Matches involving top-7 teams (2015/16):", nrow(matches_1516_top), "\n")
```

```
## Matches involving top-7 teams (2015/16): 124
```

```r
cat("Unique teams in filtered set:\n")
```

```
## Unique teams in filtered set:
```

```r
unique(c(matches_1516_top$home_team.home_team_name,
         matches_1516_top$away_team.away_team_name)) %>%
  sort() %>% print()
```

```
##  [1] "Augsburg"                 "Bayer Leverkusen"
##  [3] "Bayern Munich"            "Borussia Dortmund"
##  [5] "Borussia Mönchengladbach" "Darmstadt 98"
##  [7] "Eintracht Frankfurt"      "FC Köln"
##  [9] "FSV Mainz 05"             "Hamburger SV"
## [11] "Hannover 96"              "Hertha Berlin"
## [13] "Hoffenheim"               "Ingolstadt"
## [15] "Schalke 04"               "VfB Stuttgart"
## [17] "Werder Bremen"            "Wolfsburg"
```

```r
# Download and clean all events for those matches
cat("Downloading 2015/16 event data...\n")
```

```
## Downloading 2015/16 event data...
```

```r
invisible(capture.output({
  events_1516_raw <- free_allevents(MatchesDF = matches_1516_top)
}))

events_1516 <- allclean(events_1516_raw)

cat("Total events downloaded:", nrow(events_1516), "\n")
```

```
## Total events downloaded: 455187
```

```r
# Reconstruct running score for 2015/16
# We need goal_diff at the moment of each press event, per match.

goals_1516 <- events_1516 %>%
  filter(
    type.name == "Shot",
    shot.outcome.name == "Goal",
    !is.na(minute)
  ) %>%
  select(match_id, period, minute, second, team.name) %>%
  mutate(time_mins = minute + second / 60) %>%
  arrange(match_id, period, minute, second)

cat("Total goals in 2015/16 subset:", nrow(goals_1516), "\n")
```

```
## Total goals in 2015/16 subset: 373
```

```r
# Helper function: attach score state to a set of press events for one team
get_press_with_score <- function(press_events, goals_data, team_name) {

  press_events %>%
    select(id, match_id, minute, second,
           location.x, location.y, team.name) %>%
    mutate(time_mins = minute + second / 60) %>%
    group_by(match_id) %>%
    group_modify(~ {
      match_goals <- goals_data %>%
        filter(match_id == .y$match_id) %>%
        mutate(
          team_goal = team.name == team_name,
          opp_goal = team.name != team_name
        )

      if (nrow(match_goals) == 0) {
        .x %>% mutate(lev_goals = 0L, opp_goals = 0L, goal_diff = 0L)
      } else {
        .x %>%
          rowwise() %>%
          mutate(
            lev_goals = sum(match_goals$team_goal[match_goals$time_mins < time_mins],
                            na.rm = TRUE),
            opp_goals = sum(match_goals$opp_goal[ match_goals$time_mins < time_mins],
                            na.rm = TRUE),
            goal_diff = lev_goals - opp_goals
          ) %>%
          ungroup()
      }
    }) %>%
    ungroup()
}
```

```r
# For each top-7 team, compute press height in the 4 DiD cells:
#   early (30-70') × losing / not-losing
#   late  (>70') × losing / not-losing
#
# We use the same restricted window (30-70') as our main analysis for consistency

LATE_THRESHOLD <- 70
EARLY_LOW_BOUND <- 30

compute_team_profile <- function(team_name, events_data, goals_data) {

  # Get press events for this team
  press <- events_data %>%
    filter(
      team.name == team_name,
      type.name == "Pressure",
      !is.na(location.x),
      !is.na(location.y)
    )
```

```r
  if (nrow(press) < 50) {
    cat("  Too few press events for", team_name, "--> skipping\n")
    return(NULL)
  }

  # Attach score state
  press_scored <- get_press_with_score(press, goals_data, team_name)

  # Build DiD variables and filter to restricted window
  press_did_team <- press_scored %>%
    mutate(
      goal_diff = as.integer(goal_diff),
      is_losing = goal_diff < 0,
      is_late = minute > LATE_THRESHOLD
    ) %>%
    filter(
      (minute >= EARLY_LOW_BOUND & !is_late) | is_late
    )

  # Compute cell means
  cell_means <- press_did_team %>%
    group_by(is_late, is_losing) %>%
    summarise(
      mean_x = mean(location.x, na.rm = TRUE),
      se_x = sd(location.x,   na.rm = TRUE) / sqrt(n()),
      n = n(),
      .groups = "drop"
    ) %>%
    mutate(team = team_name)

  cell_means
}

# Run for all top-7 teams
cat("Computing press profiles for 2015/16 top teams...\n")
```

```
## Computing press profiles for 2015/16 top teams...
```

```r
profiles_1516 <- lapply(TOP_TEAMS_1516, function(team) {
  cat(" Processing:", team, "\n")
  compute_team_profile(team, events_1516, goals_1516)
}) %>%
  bind_rows()
```

```
##  Processing: Bayern Munich
##  Processing: Borussia Dortmund
##  Processing: Bayer Leverkusen
##  Processing: Borussia Mönchengladbach
##  Processing: FC Schalke 04
##   Too few press events for FC Schalke 04 --> skipping
##  Processing: 1. FSV Mainz 05
##   Too few press events for 1. FSV Mainz 05 --> skipping
##  Processing: Hertha BSC
##   Too few press events for Hertha BSC --> skipping
```

```r
cat("\nProfile computation complete.\n")
```

```
##
## Profile computation complete.
```

```r
cat("Teams successfully processed:",
    length(unique(profiles_1516$team)), "\n")
```

```
## Teams successfully processed: 4
```

```r
# Compute the two key metrics per team:
#   1. losing_late_height: mean press height when losing after 70'
#   2. press_maintenance: losing-late height minus not-losing-late height
#   3. did_estimate: the full DiD   (arithmetic version)

team_metrics_1516 <- profiles_1516 %>%
  select(team, is_late, is_losing, mean_x, n) %>%
  pivot_wider(
    names_from = c(is_late, is_losing),
    values_from = c(mean_x, n),
    names_sep = "_"
  ) %>%
  rename(
    early_notlosing = mean_x_FALSE_FALSE,
    early_losing = mean_x_FALSE_TRUE,
    late_notlosing = mean_x_TRUE_FALSE,
    late_losing = mean_x_TRUE_TRUE,
    n_early_notlosing = n_FALSE_FALSE,
    n_early_losing = n_FALSE_TRUE,
    n_late_notlosing = n_TRUE_FALSE,
    n_late_losing = n_TRUE_TRUE
  ) %>%
  mutate(
    # Core metric: how high do they press when losing late?
    losing_late_height = round(late_losing, 2),

    # Maintenance: do they press higher when losing vs not losing, late game?
    press_maintenance = round(late_losing - late_notlosing, 2),

    # DiD: incremental shift above natural time trend
    did_estimate = round(
      (late_losing - early_losing) - (late_notlosing - early_notlosing), 2
    ),

    season = "2015/16"
  )

# Add Leverkusen 2023/24
lev_2324_metrics <- tibble(
  team = "Bayer Leverkusen\n2023/24 *",
  early_notlosing = early_notlosing,
  early_losing = early_losing,
  late_notlosing = late_notlosing,
  late_losing = late_losing,
  n_early_notlosing = as.integer(sum(!press_did_restricted$is_late &
```

```r
                                        !press_did_restricted$is_losing)),
  n_early_losing = as.integer(sum(!press_did_restricted$is_late &
                                         press_did_restricted$is_losing)),
  n_late_notlosing = as.integer(sum( press_did_restricted$is_late &
                                        !press_did_restricted$is_losing)),
  n_late_losing = as.integer(sum( press_did_restricted$is_late &
                                         press_did_restricted$is_losing)),
  losing_late_height = round(late_losing, 2),
  press_maintenance = round(late_losing - late_notlosing, 2),
  did_estimate = round(did_estimate_r, 2),
  season = "2023/24"
)

# Combined table for all comparisons
all_metrics <- bind_rows(team_metrics_1516, lev_2324_metrics)

cat("--- Press Profile Metrics - Top Teams ---\n")
```

```
## --- Press Profile Metrics - Top Teams ---
```

```r
all_metrics %>%
  select(team, season, losing_late_height, press_maintenance,
         did_estimate, n_late_losing) %>%
  arrange(desc(losing_late_height)) %>%
  print()
```

```
## # A tibble: 5 x 6
##   team    season losing_late_height press_maintenance did_estimate n_late_losing
##   <chr>   <chr>               <dbl>             <dbl>        <dbl>         <int>
## 1 "Bayer~ 2023/~               69.5               8.8         1.36            91
## 2 "Bayer~ 2015/~               64.0               2.52       -0.96            22
## 3 "Borus~ 2015/~               63.7               6.32        3.71            97
## 4 "Borus~ 2015/~               58.4               4.21        3.69           372
## 5 "Bayer~ 2015/~               53.0              -3.74       -7.33           399
```

```r
# Figure 1: Ranked bar chart of losing-late press height
# Each bar = a team's mean press height when losing after 70'.

library(forcats)
library(stringr)

plot_bar <- all_metrics %>%
  filter(!is.na(losing_late_height)) %>%
  mutate(
    is_focal = team == "Bayer Leverkusen\n2023/24 *",
    team_label = str_wrap(team, width = 15),
    team_label = fct_reorder(team_label, losing_late_height)
  )

# Reference line: 2015/16 median
median_1516 <- median(
  all_metrics$losing_late_height[all_metrics$season == "2015/16"],
  na.rm = TRUE
)
```

```r
fig_bar <- ggplot(plot_bar,
       aes(x = team_label, y = losing_late_height, fill = is_focal)) +

  geom_col(width = 0.7, alpha = 0.9) +

  geom_text(aes(label = paste0(round(losing_late_height, 1), "m")),
            hjust = -0.15,
            size = 3.8,
            colour = "grey20",
            fontface = "bold") +

  # 2015/16 median reference line
  geom_hline(yintercept = median_1516,
             linetype = "dashed",
             colour = "grey40",
             linewidth = 0.9) +
  annotate("text",
           x = 0.6,
           y = median_1516 + 0.8,
           label = paste0("2015/16 median\n(", round(median_1516, 1), "m)"),
           hjust = 0,
           size = 3.2,
           colour = "grey40",
           fontface = "italic") +

  scale_fill_manual(
    values = c("FALSE" = "#7fbfcf", "TRUE" = "#d6604d"),
    guide = "none"
  ) +
  scale_y_continuous(
    limits = c(0, max(plot_bar$losing_late_height, na.rm = TRUE) + 8),
    breaks = seq(0, 120, 10)
  ) +
  coord_flip() +
  labs(
    title = "Counterpress Height When Losing After 70': Top Teams Compared",
    subtitle = paste0(
      "Mean press location.x (m) in losing-late state | ",
      "Restricted window (30-70' pre-treatment)\n",
      "Red = Leverkusen 2023/24 | Blue = 2015/16 top-7 Bundesliga teams with enough data
      available"
    ),
    x = NULL,
    y = "Mean Press Height (m from own goal line)"
  ) +
  theme_minimal(base_size = 13) +
  theme(
    plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
    plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
    panel.grid.minor = element_blank(),
    panel.grid.major.y = element_blank(),
    plot.background  = element_rect(fill = "white", color = NA)
  )
```

```r
print(fig_bar)
```

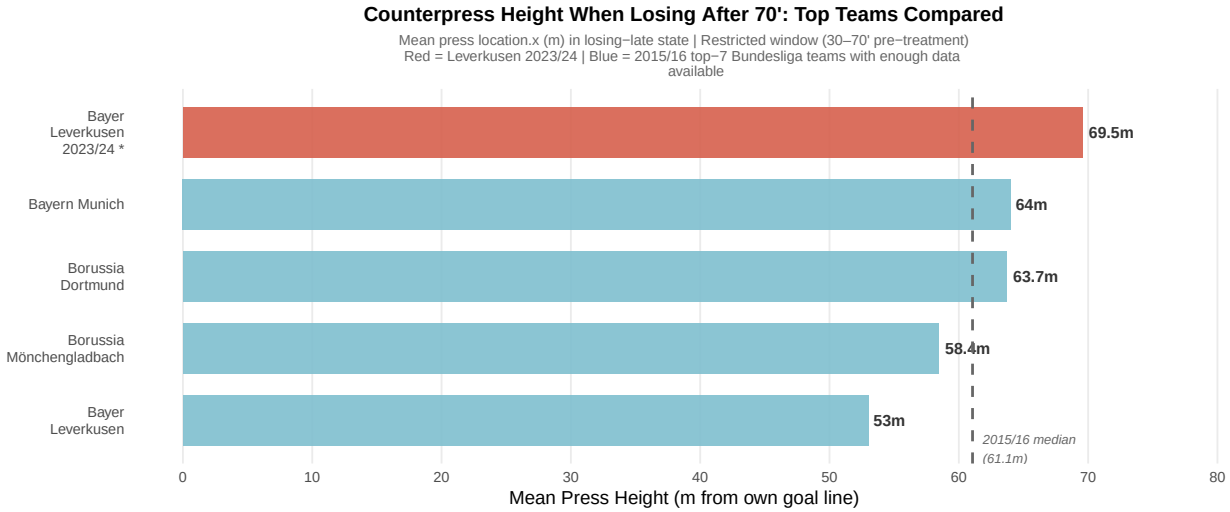


```r
ggsave("fig_benchmark_losing_late_height.png", plot = fig_bar,
       width = 10, height = 6, dpi = 300)
```

Leverkusen 2023/24 (red) sits 5.5 m above the next-best team (Bayern 64.0 m) and 8.4 m above the 2015/16 median (61.1 m). This is a large gap on a pitch that has a total length of 120 m, and it is roughly equivalent to the width of the penalty area.

Notice how Leverkusen 2015/16 (53.0 m) falls below the 2015/16 median, confirming that the aggressive losing-late press is a 2023/24 phenomenon specific to Alonso's system, not a persistent club trait. Only 4 of the 7 top teams appear, provided that Schalke, Mainz, and Hertha had insufficient pressing event data in the losing-late cell to produce reliable estimates and were excluded from the figure.

```r
# Figure 2: Scatter. Press maintenance vs losing-late height
# Two-dimensional view: x = how high they press when losing late,
# y = how much higher they press when losing vs not losing (late game).

plot_scatter <- all_metrics %>%
  filter(!is.na(losing_late_height), !is.na(press_maintenance)) %>%
  mutate(
    is_focal = team == "Bayer Leverkusen\n2023/24 *",
    team_short = case_when(
      grepl("2023/24", team)           ~ "LEV 23/24 *",
      grepl("Bayern", team)            ~ "BAY",
      grepl("Dortmund", team)          ~ "BVB",
      grepl("Leverkusen", team)        ~ "LEV 15/16",
      grepl("Gladbach|Mönchen", team)  ~ "BMG",
      grepl("Schalke", team)           ~ "S04",
      grepl("Mainz", team)             ~ "M05",
      grepl("Hertha", team)            ~ "BSC",
      TRUE ~ team
    )
  )

# Quadrant reference lines (medians of 2015/16 teams)
x_med <- median(plot_scatter$losing_late_height[plot_scatter$season == "2015/16"],
```

```r
                 na.rm = TRUE)
y_med <- median(plot_scatter$press_maintenance[ plot_scatter$season == "2015/16"],
                 na.rm = TRUE)

fig_scatter <- ggplot(plot_scatter,
       aes(x = losing_late_height, y = press_maintenance)) +

  # Quadrant shading
  annotate("rect",
           xmin = x_med, xmax = Inf, ymin = y_med, ymax = Inf,
           fill = "#d6604d", alpha = 0.06) +
  annotate("text",
           x = x_med + 0.5, y = max(plot_scatter$press_maintenance,
                                    na.rm=TRUE) * 0.92,
           label = "High press +\nMaintains under pressure",
           size = 3, colour = "#d6604d", fontface = "italic", hjust = 0) +

  # Quadrant reference lines (2015/16 medians)
  geom_vline(xintercept = x_med, linetype = "dashed",
             colour = "grey50", linewidth = 0.7) +
  geom_hline(yintercept = y_med, linetype = "dashed",
             colour = "grey50", linewidth = 0.7) +
  geom_hline(yintercept = 0, linetype = "solid",
             colour = "grey30", linewidth = 0.5) +

  # 2015/16 teams
  geom_point(data = filter(plot_scatter, !is_focal),
             colour = "#2166ac", size = 5, alpha = 0.85) +
  geom_text(data = filter(plot_scatter, !is_focal),
            aes(label = team_short),
            vjust = -0.9, size = 3.3,
            colour = "#2166ac", fontface = "bold") +

  # Leverkusen 2023/24
  geom_point(data = filter(plot_scatter, is_focal),
             colour = "#d6604d", size = 7, shape = 18) +
  geom_text(data = filter(plot_scatter, is_focal),
            aes(label = team_short),
            vjust = -1.1, size = 3.8,
            colour = "#d6604d", fontface = "bold") +

  # n labels (sample size in treatment cell)
  geom_text(aes(label = paste0("n=", n_late_losing)),
            vjust = 1.8, size = 2.8,
            colour = "grey50") +

  scale_x_continuous(breaks = seq(50, 100, 5)) +

  labs(
    title = "Pressing Under Pressure: Two-Dimensional Team Profiles",
    subtitle = paste0(
      "x = mean press height when losing after 70' (higher = more aggressive)\n",
      "y = press height gap: losing-late minus not-losing-late (positive = maintains
```

```
      press)\n",
      "Dashed lines = 2015/16 top-7 medians | n = treatment cell size"
    ),
    x = "Press Height When Losing Late (m)",
    y = "Press Maintenance Index\n(Losing Late - Not Losing Late, m)"
  ) +
  theme_minimal(base_size = 13) +
  theme(
    plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
    plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
    panel.grid.minor = element_blank(),
    plot.background = element_rect(fill = "white", color = NA)
  )

print(fig_scatter)
```

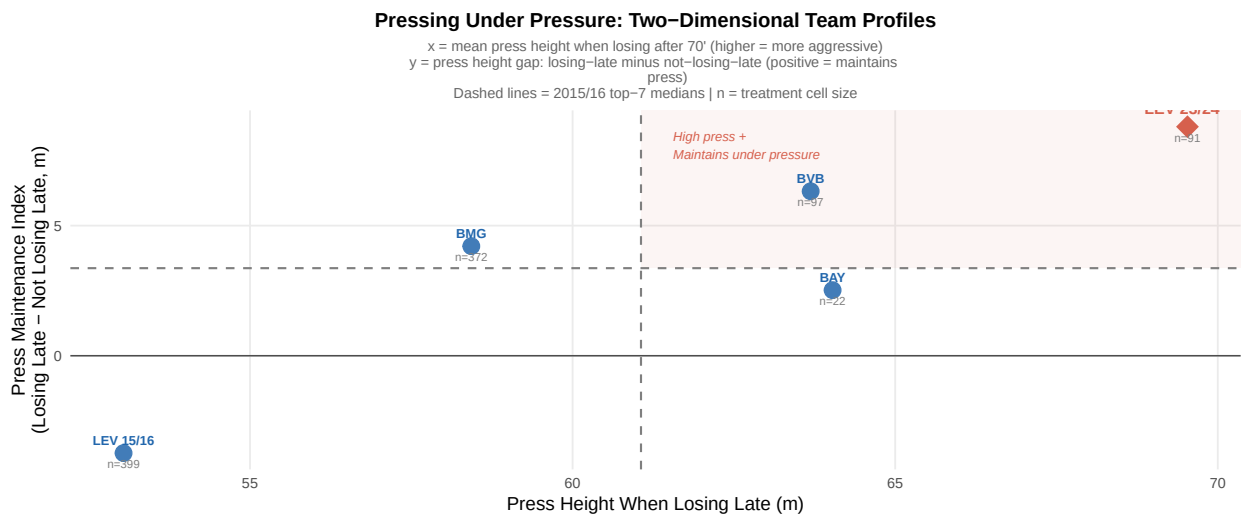


```
ggsave("fig_benchmark_scatter.png", plot = fig_scatter,
       width = 10, height = 7, dpi = 300)
```

The scatter places each team on two dimensions simultaneously: on the x-axis, raw press height when losing late (how aggressive overall) is denoted, whereas on the y-axis, press maintenance index (how much they elevate press when losing) is displayed.

Leverkusen 2023/24 (red diamond) sits in the extreme top-right quadrant, and it is the only team that combines both the highest absolute press height (69.5 m) and the strongest press maintenance (+8.80 m). No 2015/16 team occupies this quadrant with comparable values.

Dortmund (BVB) and Gladbach (BMG) show moderate maintenance but lower absolute height. Bayern Munich (BAY) presses high in absolute terms but shows relatively low maintenance (+2.52 m), suggesting their high press is a general style rather than a situational comeback response. Leverkusen 2015/16 (LEV 15/16) sits in the bottom-left, with low absolute height and negative maintenance. This is the sharpest possible contrast with their 2023/24 season.

```
# Figure 3: Press height profiles across all 4 DiD cells
# Shows each team's full pressing profile: early/late × losing/not-losing.

plot_profiles <- all_metrics %>%
  filter(!is.na(early_notlosing)) %>%
```

```r
  select(team, season, is_focal = season,
         early_notlosing, early_losing,
         late_notlosing,  late_losing) %>%
  mutate(is_focal = team == "Bayer Leverkusen\n2023/24 *") %>%
  pivot_longer(
    cols = c(early_notlosing, early_losing,
                  late_notlosing,  late_losing),
    names_to = "cell",
    values_to = "mean_x"
  ) %>%
  mutate(
    is_late = grepl("late", cell),
    is_losing = grepl("losing", cell) & !grepl("not", cell),
    period = ifelse(is_late, "Late (>70')", "Early (30-70')"),
    score = ifelse(is_losing, "Losing", "Not Losing"),
    period_n = ifelse(is_late, 1, 0),
    team_short = case_when(
      grepl("2023/24", team)          ~ "LEV\n23/24*",
      grepl("Bayern", team)           ~ "Bayern",
      grepl("Dortmund", team)         ~ "Dortmund",
      grepl("Leverkusen",team)        ~ "LEV\n15/16",
      grepl("Gladbach|Mönchen", team) ~ "Gladbach",
      grepl("Schalke", team)          ~ "Schalke",
      grepl("Mainz", team)            ~ "Mainz",
      grepl("Hertha", team)           ~ "Hertha",
      TRUE ~ team
    ),
    team_short = fct_relevel(team_short, "LEV\n23/24*", after = Inf)
  )

fig_profiles <- ggplot(plot_profiles,
       aes(x = period_n,
           y = mean_x,
           colour = score,
           group = score)) +

  geom_line(aes(linewidth = is_focal, alpha = is_focal)) +
  geom_point(aes(size = is_focal), alpha = 0.9) +

  scale_linewidth_manual(values = c("FALSE" = 0.8, "TRUE" = 1.8),
                         guide = "none") +
  scale_alpha_manual(values = c("FALSE" = 0.5, "TRUE" = 1.0),
                         guide = "none") +
  scale_size_manual(values = c("FALSE" = 2, "TRUE" = 4),
                         guide = "none") +
  scale_colour_manual(
    values = c("Losing" = "#d6604d", "Not Losing" = "#2166ac"),
    name = "Score state"
  ) +
  scale_x_continuous(
    breaks = c(0, 1),
    labels = c("Early\n(30-70')", "Late\n(>70')"),
    expand = expansion(add = c(0.2, 0.2))
```

```r
  ) +

  facet_wrap(~ team_short, ncol = 4) +

  labs(
    title = "Full DiD Cell Profiles: Leverkusen 2023/24 vs 2015/16 Top Teams",
    subtitle = paste0(
      "Each panel = one team's mean press height across the 4 DiD cells\n",
      "Bold lines = Leverkusen 2023/24 | Look for steeper Losing slope in treatment
      window"
    ),
    x = NULL,
    y = "Mean Press Height (m)"
  ) +
  theme_minimal(base_size = 12) +
  theme(
    plot.title = element_text(face = "bold", hjust = 0.5, size = 14),
    plot.subtitle = element_text(hjust = 0.5, colour = "grey40", size = 10),
    strip.text = element_text(face = "bold", size = 11),
    strip.background = element_rect(fill = "grey92", colour = NA),
    panel.grid.minor = element_blank(),
    legend.position = "top",
    plot.background = element_rect(fill = "white", color = NA)
  )

print(fig_profiles)
```

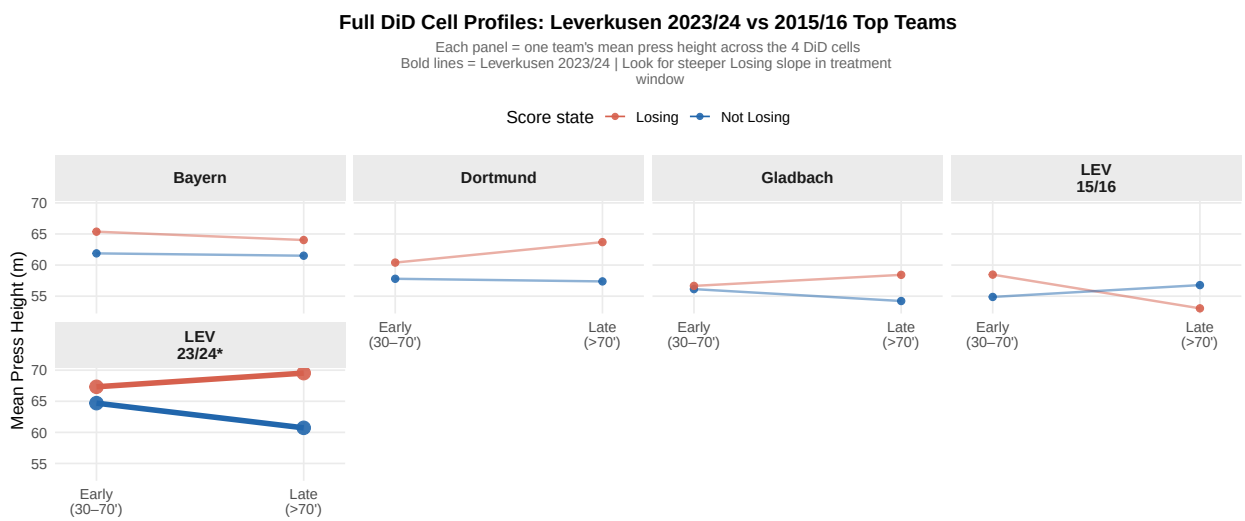


```r
ggsave("fig_benchmark_profiles.png", plot = fig_profiles,
       width = 14, height = 7, dpi = 300)
```

In the figure above, each panel shows a team's mean press height across the 4 DiD cells. Leverkusen 2023/24 (bottom-left, bold lines) clearly displays an upward slope on the Losing line (+2.5 m early to late) while Not Losing drops sharply (-4 m). This creates the widening gap that defines comeback pressing.

For the case of Bayern Munich, both lines are relatively flat and close together, which signals high general pressing intensity but little situational adjustment.

For the case of Dortmund, the losing line rises notably from ~60 m to ~64 m. This is the 2015/16 team with

the most similar profile to Leverkusen 2023/24, though at a much lower absolute level.

Gladbach's lines converge late, as the gap between losing and not-losing narrows after 70', suggesting they do not differentiate by score state.

The Leverkusen 2015/16 is perhaps the most insightful panel, since the Losing line drops from early to late while Not Losing rises, producing a crossing pattern. Under Xabi Alonso's predecessor, losing late meant pressing lower, the complete tactical opposite of the 2023/24 system.

```
# Final summary table
cat("--- Benchmark Summary Table ---\n\n")
```

```
## --- Benchmark Summary Table ---
```

```
all_metrics %>%
  filter(!is.na(losing_late_height)) %>%
  arrange(desc(losing_late_height)) %>%
  mutate(
    percentile_rank = round(
      rank(losing_late_height) / n() * 100, 0
    )
  ) %>%
  select(
    Team = team,
    Season = season,
    `Press ht (losing late)` = losing_late_height,
    `Press maintenance` = press_maintenance,
    `DiD estimate` = did_estimate,
    `N (losing late)` = n_late_losing,
    `Percentile rank` = percentile_rank
  ) %>%
  print()
```

```
## # A tibble: 5 x 7
##   Team          Season Press ht (losing lat~1 `Press maintenance` `DiD estimate`
##   <chr>         <chr>                   <dbl>               <dbl>          <dbl>
## 1 "Bayer Lever~ 2023/~                   69.5                8.8            1.36
## 2 "Bayern Muni~ 2015/~                   64.0                2.52          -0.96
## 3 "Borussia Do~ 2015/~                   63.7                6.32           3.71
## 4 "Borussia Mö~ 2015/~                   58.4                4.21           3.69
## 5 "Bayer Lever~ 2015/~                   53.0               -3.74          -7.33
## # i abbreviated name: 1: `Press ht (losing late)`
## # i 2 more variables: `N (losing late)` <int>, `Percentile rank` <dbl>
```

---

# Conclusion

Our analysis investigated whether Bayer Leverkusen's 2023/24 unbeaten season was underpinned by an anomalous spatial pressing response under adverse conditions, or if their late-game aggression was a standard among elite Bundesliga teams. The spatial and econometric evidence hint at the former: under Xabi Alonso, Leverkusen engineered a distinctive, highly scalable tactical system defined by its extreme spatial escalation when chasing a deficit. By integrating spatial KDE heatmaps, zonal raster statistics, and a Difference-in-Differences (DiD) framework, several core tenets of their identity emerge:

Leverkusen does not press blindly. Their defensive structure is designed to funnel opponents into the advanced flanks, where their pressing traps peak at a highly efficient 36% to 39% success rate. When trailing

after the 70th minute, Leverkusen breaks from natural game progression. Rather than dropping deeper as matches wear on, they enact a coordinated forward shift. While the aggregate DiD estimator (+1.36 m in the restricted 30–70 minute specification) is visually muted, conditioning the model on opponent quality reveals massive tactical escalation against mid-tier teams (+9.97 m). Against top-tier sides, they adopt a pragmatic stubbornness (+2.78 m), refusing to concede territory to dangerous counter-attacking opponents. Benchmarking this behavior against the top seven Bundesliga teams from the 2015/16 season with available data confirms that Alonso's system is an outlier, not an elite norm. Leverkusen 2023/24 posted the highest absolute press height when losing late (69.5 m) and the most aggressive press maintenance index (+8.80 m). Crucially, this is entirely a product of the current regime; the 2015/16 iteration of Leverkusen exhibited the exact opposite behavior, dropping their pressing lines significantly when behind.

We must transparently acknowledge the statistical boundaries of this study. The large p-values and wide confidence intervals in our DiD estimators fall short of traditional significance thresholds. However, in the context of sports analytics, this is a feature rather than a bug of an unbeaten season. Bayer Leverkusen simply spent so few late-game minutes in a losing state that achieving robust statistical power is mathematically impossible.

Ultimately, the spatial footprints left by this team tell a clear story. Bayer Leverkusen's historic 2023/24 campaign was not just a product of technical brilliance or luck, but of a rigorously drilled, spatially anomalous tactical system that enclosed opponents exactly when the stakes were highest.

---

# References


Andrienko, G., Andrienko, N., Budziak, G., Dykes, J., Fuchs, G., von Landesberger, T., & Weber, H. (2017). Visual analysis of pressure in football. Data Mining and Knowledge Discovery, 31(6), 1793–1839. https://doi.org/10.1007/s10618-017-0513-2

Bauer, P., & Anzer, G. (2021). Data-driven detection of counterpressing in professional football: A supervised machine learning task based on synchronized positional and event data with expert-based feature extraction. Data Mining and Knowledge Discovery, 35(5), 2009–2049. https://doi.org/10.1007/s10618-021-00763-7

Epasinghege Dona, N., & Swartz, T. B. (2024). Causal analysis of tactics in soccer: The case of throw-ins. IMA Journal of Management Mathematics, 35(1), 111–126. https://doi.org/10.1093/imaman/dpad022

Peters, A., Parmar, N., Davies, M., & James, N. (2025). A rule-based approach to classify counterpressing – analysis of its risks and relationship with rest defence. International Journal of Performance Analysis in Sport. Advance online publication. https://doi.org/10.1080/24748668.2025.2473799

Vogelbein, M., Nopp, S., & Hökelmann, A. (2014). Defensive transition in soccer–are prompt possession regains a measure of success? A quantitative analysis of German Fußball-Bundesliga 2010/2011. Journal of Sports Sciences, 32(11), 1076–1083. https://doi.org/10.1080/02640414.2013.879671